

THE KERF PROGRAMMING LANGUAGE

John Earnest

(Work In Progress)

This manual is a reference guide to Kerf, a concise multi-paradigm language with an emphasis on high-performance data processing. For the latest information and licensing inquiries, please consult:

<http://www.kerfsoftware.com>

December 31, 2015

Contents

1	Introduction	7
1.1	Background	7
1.2	Conventions	7
1.3	Using the REPL	8
1.3.1	Command-Line Arguments and Scripting	8
1.3.2	The REPL	9
1.4	Examples	10
2	Installation	11
2.1	Installing (Evaluation)	11
2.2	Installing (Building from Source)	12
2.3	Adding Kerf to your Path	12
3	Terminology	13
3.1	Atomicity	13
3.2	Combinators	14
3.3	Matrix	15
3.4	Truthy	15
3.5	Valence	15
3.6	Vector	15
4	Datatypes	16
4.1	Numbers	16
4.2	Lists and Vectors	17
4.3	Strings and Characters	18
4.4	Timestamps	19
4.5	Maps and Tables	20
4.6	Special Identifiers	21
4.7	Index and Enum	22
5	Syntax	23
5.1	Expressions	23
5.2	Indexing	25
5.3	Assignment	26
5.4	Control Structures	28
5.4.1	Conditionals	28
5.4.2	Loops	29
5.4.3	Function Declarations	30
6	SQL	31
6.1	INSERT	31
6.2	DELETE	33
6.3	SELECT	34
6.3.1	WHERE	35
6.3.2	GROUP BY	36
6.4	UPDATE	37
6.5	Joins	38
6.5.1	Left Join	39
6.5.2	Asof Join	40
6.6	Limiting	41

6.7	Ordering	41
6.8	Performance	42
7	Input/Output	43
7.1	General I/O	43
7.2	File I/O	44
7.3	Striped Files	45
7.4	Network I/O (IPC)	46
8	Foreign Function Interface	47
8.1	Overview	47
8.2	Reference Counting	49
8.2.1	Releasing	49
8.2.2	Retaining	50
8.3	Internal Representations	52
8.3.1	Integers	52
8.3.2	Floats	52
8.3.3	Characters	52
8.3.4	Stamps	52
8.3.5	Vectors	52
8.4	Attribute Flags	53
8.5	Native API	54
8.5.1	kerf_api_append - Append to List	54
8.5.2	kerf_api_call_dyad - Call Dyad	54
8.5.3	kerf_api_call_monad - Call Monad	55
8.5.4	kerf_api_call_nilad - Call Nilad	55
8.5.5	kerf_api_copy_on_write - Copy On Write	55
8.5.6	kerf_api_get - Kerf Get	56
8.5.7	kerf_api_interpret - Interpret String	56
8.5.8	kerf_api_len - Kerf Length	57
8.5.9	kerf_api_new_charvec - New Kerf String	57
8.5.10	kerf_api_new_float - New Kerf Float	57
8.5.11	kerf_api_new_int - New Kerf Integer	57
8.5.12	kerf_api_new_kerf - New Kerf Object	58
8.5.13	kerf_api_new_list - New Kerf List	58
8.5.14	kerf_api_new_map - New Kerf Map	59
8.5.15	kerf_api_new_stamp - New Kerf Timestamp	59
8.5.16	kerf_api_nil - Kerf Nil	60
8.5.17	kerf_api_release - Release Kerf Reference	60
8.5.18	kerf_api_retain - Retain Kerf Reference	60
8.5.19	kerf_api_set - Kerf Set	60
8.5.20	kerf_api_show - Show Kerf Object	61
9	Built-In Function Reference	62
9.1	abs - Absolute Value	62
9.2	acos - Arc Cosine	62
9.3	add - Add	62
9.4	and - Logical AND	62
9.5	ascend - Ascending Indices	63
9.6	asin - Arc Sine	63
9.7	asof_join - Asof Join	63
9.8	atan - Arc Tangent	64

9.9	atlas - Atlas Of	64
9.10	atom - Is Atom?	64
9.11	avg - Average	64
9.12	between - Between?	65
9.13	btree - BTree	65
9.14	bucketed - Bucket Values	65
9.15	build_table_from_csv - Build Table From CSV File	65
9.16	build_table_from_fixed_file - Build Table From Fixed-Width File	65
9.17	build_table_from_psv - Build Table From PSV File	66
9.18	build_table_from_tsv - Build Table From TSV File	66
9.19	car - Contents of Address Register	66
9.20	cdr - Contents of Data Register	66
9.21	ceil - Ceiling	66
9.22	char - Cast to Char	67
9.23	close_socket - Close Socket	67
9.24	combinations - Combinations	67
9.25	cos - Cosine	68
9.26	cosh - Hyperbolic Cosine	68
9.27	count - Count	68
9.28	count_nonnull - Count Non-Nulls	68
9.29	count_null - Count Nulls	68
9.30	cross - Cartesian Product	69
9.31	deal - Deal	69
9.32	descend - Descending Indices	70
9.33	dir_ls - Directory Listing	70
9.34	display - Display	70
9.35	distinct - Distinct Values	71
9.36	divide - Divide	71
9.37	dlload - Dynamic Library Load	72
9.38	dotp - Dot Product	72
9.39	drop - Drop Elements	72
9.40	enlist - Enlist Element	73
9.41	enum - Enumeration	73
9.42	enumerate - Enumerate Items	73
9.43	equal - Equal?	74
9.44	equals - Equals?	74
9.45	erf - Error Function	74
9.46	erfc - Complementary Error Function	74
9.47	eval - Evaluate	75
9.48	except - Except	75
9.49	exit - Exit	75
9.50	exp - Natural Exponential Function	75
9.51	explode - Explode	76
9.52	filter - Filter	76
9.53	first - First	76
9.54	flatten - Flatten	77
9.55	float - Cast to Float	77
9.56	floor - Floor	77
9.57	greater - Greater Than?	78
9.58	greatereq - Greater or Equal?	78
9.59	has_column - Table Has Column?	78

9.60	has_key - Has Key?	79
9.61	hash - Hash	79
9.62	hashed - Hashed	79
9.63	ident - Identity	80
9.64	ifnull - If Null?	80
9.65	implode - Implode	80
9.66	in - In?	80
9.67	index - Index	80
9.68	indexed - Indexed	81
9.69	int - Cast to Int	81
9.70	intersect - Set Intersection	81
9.71	isnull - Is Null?	81
9.72	join - Join	82
9.73	json_from_kerf - Convert Kerf to JSON	82
9.74	kerf_from_json - Convert JSON to Kerf	83
9.75	kerf_type - Type Code	83
9.76	kerf_type_name - Type Name	84
9.77	last - Last	84
9.78	left_join - Left Join	84
9.79	len - Length	84
9.80	less - Less Than?	85
9.81	lesseq - Less or Equal?	85
9.82	lg - Base 2 Logarithm	85
9.83	lines - Lines From File	85
9.84	ln - Natural Logarithm	86
9.85	load - Load Source	86
9.86	log - Logarithm	86
9.87	lsq - Least Squares Solution	87
9.88	map - Make Map	87
9.89	match - Match?	87
9.90	mavg - Moving Average	88
9.91	max - Maximum	88
9.92	maxes - Maximums	88
9.93	mcount - Moving Count	88
9.94	median - Median	89
9.95	meta_table - Meta Table	89
9.96	min - Minimum	89
9.97	mins - Minimums	89
9.98	minus - Minus	90
9.99	minv - Matrix Inverse	90
9.100	mmax - Moving Maximum	90
9.101	mmin - Moving Minimum	90
9.102	mmul - Matrix Multiply	91
9.103	mod - Modulus	91
9.104	msum - Moving Sum	91
9.105	negate - Negate	92
9.106	negative - Negative	92
9.107	ngram - N-Gram	92
9.108	not - Logical Not	92
9.109	noteq - Not Equal?	93
9.110	now - Current DateTime	93

9.111	now_date - Current Date	93
9.112	now_time - Current Time	93
9.113	open_socket - Open Socket	94
9.114	open_table - Open Table	94
9.115	or - Logical OR	94
9.116	order - Order	94
9.117	out - Output	94
9.118	part - Partition	95
9.119	permutations - Permutations	95
9.120	plus - Plus	95
9.121	pow - Exponentiation	96
9.122	powerset - Power Set	96
9.123	rand - Random Numbers	96
9.124	range - Range	97
9.125	read_from_path - Read From Path	97
9.126	read_stripped_from_path - Read Striped File From Path	98
9.127	read_table_from_csv - Read Table From CSV File	98
9.128	read_table_from_delimited_file - Read Table From Delimited File	98
9.129	read_table_from_fixed_file - Read Table From Fixed-Width File	99
9.130	read_table_from_tsv - Read Table From TSV File	99
9.131	rep - Output Representation	100
9.132	repeat - Repeat	100
9.133	reserved - Reserved Names	100
9.134	reverse - Reverse	101
9.135	rsum - Running Sum	101
9.136	run - Run	101
9.137	search - Search	102
9.138	seed_prng - Set random seed	102
9.139	send_async - Send Asynchronous	102
9.140	send_sync - Send Synchronous	102
9.141	setminus - Set Disjunction	102
9.142	shell - Shell Command	102
9.143	shift - Shift	103
9.144	shuffle - Shuffle	103
9.145	sin - Sine	103
9.146	sinh - Hyperbolic Sine	104
9.147	sleep - Sleep	104
9.148	sort - Sort	104
9.149	sort_debug - Sort Debug	105
9.150	split - Split List	105
9.151	sqrt - Square Root	105
9.152	stamp_diff - Timestamp Difference	106
9.153	std - Standard Deviation	106
9.154	string - Cast to String	106
9.155	subtract - Subtract	106
9.156	sum - Sum	106
9.157	tables - Tables	107
9.158	take - Take	107
9.159	tan - Tangent	107
9.160	tanh - Hyperbolic Tangent	107
9.161	times - Multiplication	108

9.162	timing - Timing	108
9.163	tolower - To Lowercase	108
9.164	toupper - To Uppercase	108
9.165	transpose - Transpose	109
9.166	trim - Trim	109
9.167	type_null - Type Null	109
9.168	uneval - Uneval	110
9.169	union - Set Union	110
9.170	unique - Unique Elements	110
9.171	var - Variance	110
9.172	which - Which	111
9.173	write_csv_from_table - Write CSV From Table	111
9.174	write_delimited_file_from_table - Write Delimited File From Table	111
9.175	write_striped_to_path - Write Striped File To Path	111
9.176	write_text - Write Text	112
9.177	write_to_path - Write to Path	112
9.178	xbar - XBar	112
9.179	xkeys - Object Keys	112
9.180	xvals - Object Values	112
10 Combinator Reference		113
10.1	converge - Converge	113
10.2	deconverge - Deconverge	113
10.3	fold - Fold	114
10.4	mapback - Map Back	114
10.5	mapcores - Map to Cores	115
10.6	mapdown - Map Down	116
10.7	mapleft - Map Left	116
10.8	mapright - Map Right	116
10.9	reconverge - Reconverge	117
10.10	reduce - Reduce	117
10.11	refold - Refold	117
10.12	rereduce - Re-Reduce	117
10.13	unfold - Unfold	117
11 Global Reference		118
11.1	.Argv - Arguments	118
11.2	Math	118
11.2.1	.Math.BILLION - Billion	118
11.2.2	.Math.E - E	118
11.2.3	.Math.TAU - Tau	118
11.3	Net	118
11.3.1	.Net.client - Client	118
11.3.2	.Net.on_close - On Close	118
11.4	Parse	119
11.4.1	.Parse.strptime_format - Time Stamp Format	119
11.4.2	.Parse.strptime_format2 - Time Format	119
12 Programming Techniques		120
12.1	Reversing a Map	120
12.2	Run-Length Encoding	123
12.3	HTTP Fetching	127

1 Introduction

Kerf is a programming language built on pragmatism, borrowing ideas from many popular tools. The syntax of Kerf will be familiar enough to anyone who has programmed in C, Python or VBA. Data is described using syntax from JSON (JavaScript Object Notation), a text-based data interchange format. Queries to search, sort and aggregate data can be performed using SQL syntax. Kerf's built-in commands have aliases which allow programmers to use names and terms they are already used to.

Beneath this friendly syntax, Kerf exposes powerful ideas inspired by the language APL and its descendants. APL has a well-earned reputation for extreme concision, and with practice you will find that Kerf similarly permits you to say a great deal with a few short words. Coming from other programming languages, you may be surprised by how much you can accomplish without writing loops, using conditional statements or declaring variables. Kerf provides a fluid interface between your intentions and your data.

1.1 Background

The Kerf team was first introduced to array languages in 2006 by Dennis Shasha at NYU. This led to work with Arthur Whitney's & Kx System's kdb+ family of languages at the investment banks Cantor Fitzgerald and Merrill Lynch. A precursor language to Kerf called Kona was started around 2009. Kona was open-sourced in Summer 2010 and improved upon with the community for the following four years. Lessons from Kona would inspire Kerf.

The first lines of code for Kerf were written in Summer 2014. Kerf officially launched as a product in Spring 2015. Kerf is the team's third major programming language release, and arguably a fifth generation language and database system. It draws on lessons from over twenty years of programming. It is mature technology.

The name "Kerf" comes from a term in woodworking- the cut made by a saw. It springs from the Old English "cyrf", the action of cutting. It's short, strong, and simple.

1.2 Conventions

Throughout this manual, the names of functions and commands will be shown in a monospaced font. Transcripts of terminal sessions will be shown with sections typed by the user in blue:

```
KerF> range 6
[0, 1, 2, 3, 4, 5]
KerF> sum(5, range 6)
20
```

Sometimes examples will contain comments, colored orange to help set them apart from code:

```
2+3;    // kerf uses c-style line comments
```


1.3 Using the REPL

A Read-Evaluate-Print Loop (REPL) is an interactive console session that allows you to type code and see results. The REPL is the main way you will be interacting with Kerf. If Kerf is in the current directory, you can start the REPL by typing `./kerf` and pressing return, and if you have installed Kerf in your path, you can simply type `kerf`. The rest of this discussion will assume the latter case.

1.3.1 Command-Line Arguments and Scripting

Kerf accepts several command-line flags to control its behavior. Throughout this manual we will be using the `-q` flag for some examples to avoid showing the Kerf startup logo for the sake of brevity.

Flag	Arguments	Behavior
<code>-q</code>	-	Suppress the startup banner.
<code>-l</code>	-	Enable debug logging.
<code>-e</code>	String Expression	Execute an expression.
<code>-x</code>	String Expression	Evaluate an expression and print the result.
<code>-p</code>	Port Number	Specify a listening port for starting an IPC server.

Summary of Command-Line Flags

The `-e` and `-x` flags differ by whether or not they display the result of a calculation. Either will exit the interpreter when complete:

```
> kerf -x "2+3"
5
> kerf -e "2+3"
>
```

If you provide a filenames as command-line arguments, the contents of those files will be executed before opening the REPL. You may wish to conclude scripts with `exit(0)` so that they execute and then self-terminate:

```
> cat example.kerf
display join unfold range(10)
exit(0)
> kerf example.kerf
[0,
[0, 1],
[0, 1, 2],
[0, 1, 2, 3],
[0, 1, 2, 3, 4],
[0, 1, 2, 3, 4, 5],
[0, 1, 2, 3, 4, 5, 6],
[0, 1, 2, 3, 4, 5, 6, 7],
[0, 1, 2, 3, 4, 5, 6, 7, 8],
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]]
>
```

You could also accomplish the same by using `-x` and `load`:

```
> kerf -x "load 'example.kerf'"
```

1.3.2 The REPL

The REPL always begins with the `KeRF>` prompt. Type an expression, press return and the result will be printed, followed by an empty line. If an expression returns null, it will appear as an empty line. Trailing whitespace will generally be elided from REPL transcripts in this manual.

```
KeRF> 1+3 5 7
      [4, 6, 8]

KeRF> null

KeRF>
```

If you type several expressions separated by semicolons (;), each will be executed left to right and the value returned by the final expression will be printed. An empty expression returns null, so if you end a statement with a semicolon it will effectively suppress printing the result. Transcripts in this manual will often use this technique for the sake of brevity.

```
KeRF> one: 1; two: 2
      2
KeRF> one
      1
KeRF> a: 3 5 7;
KeRF>
```

If you type an expression with an unbalanced number of [, (or {, the REPL will prompt you with > to complete the expression on the next line. Remember, newlines and semicolons are always equivalent:

```
KeRF> [1 2 3
> 4 5 6]
      [[1, 2, 3],
       [4, 5, 6]]
```

1.4 Examples

Let's look at a few short Kerf snippets to get a taste of the language:

Are two strings anagrams?

```
def are_anagrams(a, b) {  
    return (sort a) match (sort b)  
}
```

```
KeRF> are_anagrams("baton", "stick")  
0  
KeRF> are_anagrams("setecastronomy", "toomanysecrets")  
1
```

Gather simple statistics for random data using SQL syntax:

```
KeRF> data: {{a: rand(100, 5)}};  
KeRF> SELECT count(a) AS items, avg(a) AS average FROM data
```

items	average
100	1.82

Load a text file and interactively query it:

```
KeRF> characters: tolower flatten lines "flour.txt"  
"flour is a powder made by grinding uncooked cereal grains or other seeds or roots (like  
cassava). it is the main ingredient of bread, which is a staple food for many cultures,  
making the availability o..."  
KeRF> sum characters in "aeiou"  
182  
KeRF> count characters  
540  
KeRF> 5 take ` " " explode characters  
["flour", "is", "a", "powder", "made"]
```

Iteratively calculate terms of the Fibonacci sequence without using explicit loops:

```
KeRF> 6 {[x] last(x) join sum x} deconverge 1 1  
[[1, 1], [1, 2], [2, 3], [3, 5], [5, 8], [8, 13], [13, 21]]  
KeRF> first mapdown 10 {[x] last(x) join sum x} deconverge 1 1  
[1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89]
```

2 Installation

Pre-compiled evaluation copies of Kerf for Linux and OSX can be obtained from the project's public page on GitHub. Currently, evaluation copies expire periodically. Kerf is gaining new features frequently, so make sure you keep your installation up to date.

<https://github.com/kevinlawler/kerf>

2.1 Installing (Evaluation)

Fetch a local copy of the Kerf repository using `git clone`:

```
/Users/john/Desktop> git clone https://github.com/kevinlawler/kerf.git
Cloning into 'kerf'...
remote: Counting objects: 538, done.
remote: Total 538 (delta 0), reused 0 (delta 0), pack-reused 538
Receiving objects: 100% (538/538), 29.26 MiB | 1.65 MiB/s, done.
Resolving deltas: 100% (225/225), done.
Checking connectivity... done
/Users/john/Desktop> cd kerf/
/Users/john/Desktop/kerf>
```

The binaries located in the `/KerfREPL/linux` and `/KerfREPL/osx` directories are for 64-bit Linux and OSX, respectively. They are statically linked and include all library dependencies. Installation is as simple as placing the binary in a desired directory. Let's place it in a directory called `/opt/kerf` so that it is accessible by all users. It will be necessary to use the `sudo` command when creating this directory, as the base directory is owned by root:

```
/Users/john/Desktop> sudo mkdir /opt/kerf
Password:
/Users/john/Desktop> sudo cp KerfREPL/osx/kerf /opt/kerf/
/Users/john/Desktop> cd /opt/kerf
/opt/kerf> ls
kerf
```

You can then invoke it from the command line. The `-q` option (*quiet*) suppresses the Kerf logo at startup, for the purposes of brevity in these transcripts.

```
/opt/kerf> ./kerf -q
KeRF> 2+3
5
KeRF> exit(0)
/opt/kerf>
```

From time to time, you should use the command `git pull` from within the directory you created via `git clone` to fetch the latest build of Kerf from this repository. Then simply repeat the above steps to move the fresh binary into its normal resting place.

2.2 Installing (Building from Source)

If you have been granted access to the Kerf source code, you can build your own binaries. From the base source directory, invoke `make clean` to remove any temporary or compiled files and then `make` to build a fresh set of binaries for your OS:

```
/Users/john/Desktop/kerf-source> make clean
find manual/ -type f -not -name '*.tex' | xargs rm
rm -f -r kerf kerf_test ./obj/*.o
/Users/john/Desktop/kerf-source> make
clang -rdynamic -m64 -w -Os -c alter.c -o obj/alter.o

...

/Users/john/Desktop/kerf-source>
```

This process will produce a `kerf` executable in the source directory. You can then follow the steps described in the above section to place this in a directory accessible by other users or simply run it in place. Compiling from source will also produce a binary named `kerf_test` which executes self-tests at startup and executes in debug mode, printing extra information when errors are encountered. The debug mode binary is particularly helpful when debugging dynamic libraries, as it can detect many types of memory leak at shutdown.

2.3 Adding Kerf to your Path

You may wish to add the Kerf binary to your `PATH` so that it can be accessed more easily. If you've placed the binary in the directory `/opt/kerf/`, edit `~/.profile` and add a line to initialize this setting.

If you're using `bash` (The default shell on OSX):

```
export PATH=/opt/kerf/:$PATH
```

If you're using `tcsh` or `csch` (The default shell in some Linux distros):

```
set path = ($path /opt/kerf)
```

Open a fresh terminal or type `source ~/.profile` and you should now be able to invoke the `kerf` command from any directory.

3 Terminology

Kerf uses terminology from databases, statistics and array-oriented programming languages like APL. This section will serve as a primer for concepts which may seem unfamiliar.

3.1 Atomicity

Atomicity describes the manner in which values are *conformed* by particular functions.

A function which is not atomic will simply be applied to its arguments, behaving the same whether they are lists or atoms. `enlist` is not atomic:

```
KerF> enlist 4
[4]
KerF> enlist [1, 9, 8]
[[1, 9, 8]]
```

A unary function which is *atomic* will completely decompose any nested lists in the argument, operate on each atom separately, and then reassemble these results to match the shape of the original argument. Another way to think of this is that atomic functions “penetrate” to the atoms of their arguments. `not` is atomic:

```
KerF> not 1
0
KerF> not [1, 0, 0, 1, 0]
[0, 1, 1, 0, 1]
KerF> not [1, 0, [1, 0], [0, 1], 0]
[0, 1, [0, 1], [1, 0], 1]
```

Things get more interesting when dealing with *fully atomic* binary functions. The shapes of the arguments do not have to be identical, but they must recursively *conform*. Atoms conform with atoms. Lists conform with atoms and vice versa. Lists only conform with other lists if their lengths match and each successive pairing of their elements conforms. `add` is fully atomic:

```
KerF> add(1, 2)
3
KerF> add(1 3 5, 10)
[11, 13, 15]
KerF> add(1 2 3, 4 5 6)
[5, 7, 9]

KerF> add(1 2 3, 4 5)
add(1 2 3, 4 5)
^
Length error
```

3.2 Combinators

In the context of Kerf, a *combinator* is an operator which controls how a *function* is applied to *values*. Combinators express abstract patterns which recur frequently in programming. For example, consider the following loop:

```
function mysum(a) {  
    s: 0  
    for(i: 0; i < len(a); i: i+1) { s: add(s, a[i]) }  
    return s  
}
```

We’re iterating over the indices of the list `x` from left to right, accumulating a result into the variable `s`. On each iteration, we take the previous `s` and combine it with the current element of `a` via the function `add`. This pattern is captured by the combinator `fold`, which takes a function as a left argument and a list as a right argument:

```
Kerf> add fold 37 15 4 8  
64
```

Think of `fold` as applying its function argument between the elements of its list argument:

```
Kerf> ((37 add 15) add 4) add 8  
64
```

In this particular case we could have simply used the built-in function `sum`, but `fold` can be applied to any function- including functions you define yourself. Combinators are generally much more concise than writing explicit loops, and by virtue of having fewer “moving parts” avoid many classes of potential mistake entirely. If you use `fold` it isn’t possible to have an “off-by-one” index error accessing elements of the list argument and several useful base cases are handled automatically. Familiarize yourself with all of Kerf’s combinators- with practice, you may find you hardly ever need to use `for`, `do` and `while` loops at all!

Kerf understands the patterns combinators express and can sometimes perform dramatic optimizations when they are used in particular combinations with built-in functions or data with specific properties:

```
Kerf> timing(1);  
Kerf> max fold range 50000  
49999  
0 ms  
Kerf> {[a,b] max(a, b)} fold range 50000  
49999  
3.2 s
```

The former example allows Kerf to recognize the opportunity for short-circuiting `max fold` because the result of `range` is sorted. When folding a user-declared lambda, it must construct and then reduce the entire list.

3.3 Matrix

A *matrix* is a *vector* of *vectors* of uniform length and type. For example, the following is a matrix:

```
[[1, 2, 3],  
 [4, 5, 6],  
 [7, 8, 9]]
```

But this is not a matrix, because the rows are not of uniform length:

```
[[1, 2],  
 [3, 4, 5, 6],  
 [7, 8, 9]]
```

3.4 Truthy

Kerf does not have a special “boolean” type for representing the values *true* and *false*. By convention, the values 1 and 0 are used in most cases- this is particularly helpful in combination with the *which* operator:

```
Kerf> [true, false, true, true]  
      [1, 0, 1, 1]  
Kerf> which 1 0 1 1  
      [0, 2, 3]
```

In some situations, Kerf permits a broader range of values to behave like true, or *truthy*. Values which behave like false are naturally *falsey*. Any numeric zero or null is falsey, and any other value is considered *truthy*.

```
Kerf> if (-5) { display "yep" }  
"yep"  
Kerf> if ([]) { display "yep" }  
"yep"  
Kerf> if ({} ) { display "yep" }  
"yep"  
Kerf> if (NaN) { display "yep" }  
"yep"  
Kerf> if (0) { display "nope" }  
Kerf> if (0.0) { display "nope" }  
Kerf> if (null) { display "nope" }
```

Heavy reliance on truthiness can lead to very confusing code. Prefer 1 and 0, or the literals *true* and *false*, whenever possible.

3.5 Valence

A function’s *valence* is the number of arguments it takes. For example, *add* is a *binary function* which takes two arguments and thus has a valence of 2. *not*, on the other hand, is a *unary function* which takes a single argument and thus has a valence of 1. The term draws an analogy to linguistics and in turn chemistry, describing the way words and molecules form compounds.

3.6 Vector

A *vector* is a list of elements with a uniform type. If a list contains more than one type of element it is sometimes referred to as a *mixed-type list*. Vectors can store data more densely than mixed-type lists and as a result are often more efficient.

4 Datatypes

4.1 Numbers

Kerf has two numeric types: *integers* (or “ints”) and *floating-point numbers* (or “floats”). The results of numeric operations between ints and floats will coerce to floats, and some operators always yield floating-point results.

Integers consist of a sequence of digits, optionally preceded by + or -. They are internally represented as 64-bit signed integers and thus have a range of $-(2^{63})$ to $2^{63} - 1$.

$$\begin{array}{r} 42 \\ 010 \\ +976 \\ -9000 \end{array}$$

Integers can additionally be one of the special values `INF`, `-INF` or `NAN`, used for capturing arithmetic overflow and invalid elements of an integer vector:

```
KerF> b: 10000000000000000000000000000000000000000000000000000  
INF  
KerF> kerf_type_name b  
"integer"  
KerF> 1 2 3[-5]  
NAN
```

Floats consist of a sequence of digits with an optional sign, decimal part and exponent. They are based on IEEE-754 double-precision 64-bit floats, and thus have a range of roughly $1.7 * 10^{\pm 308}$.

```
1.0
-379.8
-.2
.117e43
```

Floats can additionally be one of the special values `nan`, `-nan`, `inf` or `-inf`. `nan` has unusual properties in Kerf compared to most other languages. The behavior is intended to permit invalid results to propagate across calculations without disrupting other valid calculations:

```
KeRF> nan == nan
1
KeRF> nan == -nan
1
KeRF> 5/0
inf
KeRF> 0/0
nan
KeRF> -1/0
-inf
```

4.2 Lists and Vectors

Lists are ordered containers of heterogenous elements. Lists have several literal forms. A sequence of numbers separated by whitespace is a valid list. This is an “APL-style” literal:

```
1 2 3 4
47 49
```

Alternatively, separate elements with commas (,) or semicolons (;) and enclose the list in square brackets for a more explicit “JSON-style” literal:

```
[1, 2, 3]
[4;5;6]
[42]
```

This manual will use both styles throughout the examples. Naturally, these styles can be nested together. The following examples are equivalent:

```
[1 2,3 4,5,6]
[[1, 2], [3, 4], 5, 6]
[[1;2], [3;4], 5, 6]
```

If a list consists entirely of items of the same type, it is a *vector*. Vectors can be represented more compactly than mixed-type lists and thus are more cache-friendly and provide better performance. Kerf has special optimizations for vectors of timestamps, characters, floats and integers.

Vector and list types each have their own special symbol for emptiness:

```
KeRF> 0 take 1 2 3
      INT[]
KeRF> 0 take 1.0 2 3
      FLOAT[]
KeRF> 0 take "ABC"
      ""
KeRF> 0 take [now(),now()]
      STAMP[]
KeRF> 0 take [1,"A"]
      []
```

4.3 Strings and Characters

Kerf character literals begin with a backtick (``) followed by double-quotes (") or single-quotes (') surrounding the character. Character literals which do not contain escape sequences may omit the quotation marks.

```
`A
`'B'
`"C"
```

Kerf supports all JSON string escape sequences, and additionally an escape for single-quotes:

```
`"\\"      // double quote
`'\\"      // single quote
`"\\\"     // reverse solidus
`"\/"      // solidus
`"\\b"     // backspace
`"\\f"     // formfeed
`"\\n"     // newline
`"\\r"     // carriage return
`"\\t"     // horizontal tab
`"\\u0043" // 4-digit unicode literal
```

The built in function `char` can convert a number into an equivalent character:

```
KerF> char 65
`"A"
```

Strings are lists of characters, and qualify as vectors. Strings are simply enclosed in double-quotes (") or single-quotes (').

```
"Hello, World!"
'goodbye, \ncruel world...'
```

Note that a single character in quotation marks is a *string*, not a character literal:

```
KerF> kerf_type_name "A"
"character vector"
KerF> "A" match `"A"
0
KerF> first "A"
`"A"
```

When dealing with characters, spaces are considered null values:

```
KerF> isnull ` " "
1
KerF> (not isnull) filter "Text with whitespace! "
"Textwithwhitespace!"
```

4.4 Timestamps

Timestamps, or simply “stamps”, are a flexible datatype which can represent dates, times, or a complete date-time. Times are internally represented in UTC at **nanosecond granularity**. Timestamp literals use a format similar to ISO-8601 except periods (.) are used as date field separators instead of dashes (-).

```
1997.07.16           // date only
19:20:30             // time only
19:20:30.123         // time with milliseconds
1997.07.16T19:20:30  // datetime
1997.07.16T19:20:30.123 // datetime with milliseconds
```

Timestamps can be compared using the same operators as numeric types. The special builtin `stamp_diff` should be used for calculating the interval between timestamps:

```
KeRF> 2015.04.02 < 2015.05.01
1
KeRF> stamp_diff(2015.05.03, 2015.05.01)
1728000000000000
```

KeRF provides special literals for relative date-times:

```
1y      // years
10m     // months
3d      // days
1h      // hours
2i      // minutes
1s      // seconds
```

These can also be combined as a single unit. For example,

```
KeRF> 2015.01.01 + 2m + 1d
2015.03.02
KeRF> 2015.01.01 + 2m1d
2015.03.02
KeRF> 2015.01.01 - 1h1i1s
2014.12.31T22:58:59.000
```

Indexing is overloaded for timestamps to permit easy extraction of fields. The date and time fields produce a stamp and other fields produce an integer:

```
KeRF> d: now();
KeRF> d[["date", "time"]]
[2015.11.14, 23:11:18.154]
KeRF> d[["year", "month", "week", "day", "hour", "minute"]]
[2015, 11, 46, 14, 23, 11]
KeRF> d[["second", "millisecond", "nanosecond"]]
[18, 154, 154858000]
```

Timestamps are currently valid through 2262.04.11, at which point we’ll already have transitioned to 128-bit operating systems. Be sure to download the latest KeRF release when nearing the “Year 2.262k Problem”.

4.5 Maps and Tables

A *Map* is an associative data structure which binds *keys* to *values*. Kerf maps are a generalization of JSON objects. Map literals are enclosed in a pair of curly braces (`{` and `}`), and contain a series of comma delimited key-value pairs separated by a colon (`:`). JSON object syntax requires keys to be enclosed in double-quotes, but Kerf maps permit using single-quotes or bare identifiers. In any of these cases, the keys of the resulting map will be strings.

```
{ "a": 10, "b": 20 }    // JSON style
{ 'a': 10, 'b': 20 }    // optional single-quotes
{ a: 10, b: 20 }       // bare identifiers
```

A *Table* is a map in which each value is a list of equal length. Tables are enclosed in two pairs of curly braces (`{{` and `}}`) and otherwise syntactically resemble maps. If non-list values are provided, they will be wrapped in lists. Values can also be omitted entirely to produce a table with empty columns.

```
{{a: 1 2, b: 3 4}}    // columns contain [1,2] and [3,4]
{{a: 1, b: 2}}         // columns contain [1] and [2]
{{a, b}}               // columns contain [] and []
```

If tables are serialized to JSON, Kerf will insert a key `is_json_table`- this permits tables to survive round-trip conversion:

```
Kerf> json_from_kerf {{a: 1 2, b: 3 4}}
"{ \"a\": [1,2], \"b\": [3,4], \"is_json_table\": [1] }"
```

The builtins `xkeys` and `xvals` can be used to extract a list of keys or values from a map or table. The builtin `map` produces a map from a list of keys and a list of values.

```
Kerf> xkeys {a: 10 11 12, b: 20 21}
["a", "b"]
Kerf> xkeys {{a:1, b:2}}
["a", "b"]
Kerf> xvals {a: 10 11 12, b: 20 21}
[[10, 11, 12], [20, 21]]
Kerf> xvals {{a:1, b:2}}
[[1], [2]]
Kerf> map(["a","b"], [37, 99])
{a:37, b:99}
```

Many primitive operations penetrate to the values of maps and tables:

```
Kerf> 3 + {a: 10, b: 20}
a:13, b:23
Kerf> 3 + {{a: 10, b: 20}}
```

a	b
13	23

4.6 Special Identifiers

Kerf uses reserved words to identify a number of special values.

The words `true` and `false` are boolean literals, equivalent to 1 and 0, respectively:

```
KerF> true
1
KerF> false
0
```

Inside a function, the words `self` or `this` may be used to refer to the current function. This is particularly useful for performing recursive calls in anonymous “lambda” functions:

```
KerF> def foo(x) { return [this, self, x] }
  {[x] return [this, self, x]}
KerF> foo(9)
[[[x] return [this, self, x]], {[x] return [this, self, x]}, 9]
```

The words `nil` or `null` may be used interchangeably to refer to a null value:

```
KerF> nil

KerF> null

KerF> kerf_type_name nil
"null"
KerF> kerf_type_name null
"null"
```

The word `root` is a reference to Kerf’s global scope. It contains all the variables which have been defined or referenced:

```
KerF> root
{}
KerF> a: 24
24
KerF> b
Undefined token error
KerF> root
{a:24}
```

4.7 Index and Enum

Indexes and Enumerations are special lists which perform internal bookkeeping to improve the performance of certain operations.

An *Enumeration* performs *interning*. It keeps only one reference to each object and stores appearances as fixed-width indices. It is useful for storing repetitions of strings and lists, which cannot otherwise efficiently be stored as vectors. In all other respects an Enumeration appears to be a list. To create an Enumeration, use `hashed` or the unary `#` operator:

```
KeRF> a: hashed ["cherry", "peach", "cherry"]
#["cherry", "peach", "cherry"]
KeRF> kerf_type_name a
"enum"
```

Not only do Enumerations reduce the memory footprint of lists with a large number of repeated elements, they permit dramatically faster sorting. There is no benefit to making an Enumeration out of a vector of integers or floats.

```
KeRF> samples: {[x] rand(x, ["cherry", "peach", "zucchini"])};
KeRF> s1: samples(1000);
KeRF> s2: samples(10000);
KeRF> s3: samples(100000);
KeRF> timing 1
1
KeRF> sort s1;
2 ms
KeRF> sort s2;
15 ms
KeRF> sort s3;
134 ms
KeRF> sort hashed s1;
0 ms
KeRF> sort hashed s2;
4 ms
KeRF> sort hashed s3;
28 ms
```

An *Index* is a list augmented with a B-Tree. This permits more efficient lookups and range queries. Do not use an index for data that will always be sorted in ascending order- Kerf tracks sorted lists internally. To create an Index, use `indexed` or the unary `=` operator:

```
KeRF> b: indexed 3 9 0 7
=[3, 9, 0, 7]
KeRF> kerf_type_name b
"sort"
```

5 Syntax

5.1 Expressions

Calling (or *applying*) a function in Kerf resembles most conventional languages- use the name of the function followed by a parenthesized, comma-separated list of arguments:

```
KerF> add(1, 2)
3
```

If a function takes exactly one argument, the parentheses are optional. We call this style “prefix” function application:

```
KerF> negate(3)
-3
KerF> negate 3
-3
```

This syntax makes it easy to “chain” together a series of unary functions:

```
KerF> last sort unique "ALPHABETICAL"
~ "T"
KerF> last(sort(unique("ALPHABETICAL")))
~ "T"
```

If a function takes no arguments, you must remember to include parentheses- otherwise the function will be returned as a value instead of called:

```
KerF> exit
exit
KerF> exit()
>
```

If a function takes exactly two arguments, it can be placed between the first and second argument as an “infix” operator:

```
KerF> 1 add 2
3
```

Many of the most frequently used functions have symbolic aliases. For example, + can be used instead of add and * can be used instead of times. There is no functional difference between the spelled-out names for these functions and the symbols.

```
KerF> 3 * 5
15
KerF> *(3, 5)
5
KerF> 3 + 5
8
KerF> +(3, 5)
5
```


Operators have uniform precedence in Kerf. Expressions are evaluated strictly from right to left unless explicitly grouped with parentheses:

```
KerF> 3 * 4 + 1
15
KerF> 4 + 1 * 3
7
KerF> (4 + 1) * 3
15
```

Kerf's flexible syntax often provides many alternatives for writing the same expression. Select the arrangement that you feel is most clear. Adding parentheses to confusing-seeming expressions never hurts!

```
KerF> 0.5 * 3**2
4.5
KerF> times(1/2, 3**2)
4.5
KerF> divide(1, 2) * exp(3, 2)
4.5
KerF> ((1 / 2) * (3 ** 2))
4.5
```

Symbol	Unary Function	Binary Function
-	negate	minus
+	transpose	add
*	first	times
/	reverse	divide
	len (length)	maxes/or
^	enumerate	take
%	distinct	mod (modulus)
&	part (partition)	mins/and
?	which	rand (random)
#	hashed	join
!	not	map (make map)
~	atom (is atom?)	match
=	indexed	equals
<	ascend	less
>	descend	greater
\		lsq
<=		lesseq
>=		greatereq
==		equals
!=		noteq
<>		noteq
**		exp
.	floor	
.	eval (evaluate)	
:	ident (identity)	

Symbolic Aliases of Built-in Functions

5.2 Indexing

Kerf has a uniform syntax for accessing elements of lists, maps and tables. Use square brackets to the right of a variable name or expression with an index or key to look up:

```
KerF> 3 7 15[1]
7
KerF> {a: 24, b: 29}["a"]
24
```

If the provided index or key does not exist, indexing will return an appropriate type-specific null value, as provided by the `type_null` built-in function:

```
KerF> 3 7 15[9]
NAN
KerF> 3 7 15[-1]
NAN
KerF> "ABC"[9]
~" "
```

Floating-point indices to lists will be truncated, for convenience:

```
KerF> 3 7 15[0.5]
3
KerF> 3 7 15[1.6]
7
```

Indexing is right-atomic. If the indices are a list, the indexing operation will accumulate a list of results:

```
KerF> "ABC"[2 1 0 0 2 3 0 0 1]
"CBAAC AAB"
KerF> 34 19 55 32[0 1 0 1 2 2]
[34, 19, 34, 19, 55, 55]
```

One application of this type of collective indexing is the basis of `sort`:

```
KerF> a: 27 15 9 55 0
[27, 15, 9, 55, 0]
KerF> a[ascend a]
[0, 9, 15, 27, 55]
```

The shape of the result of indexing will always match the shape of the indices. Consider this example, where we index a list with a 2x2 matrix and get back a 2x2 matrix:

```
KerF> 11 22 33 44[[0 1, 2 3]]
[[11, 22],
 [33, 44]]
```

Indexing a particular element from a multidimensional structure requires several indexing operations:

```
KerF> [11 22, 33 44][0][1]
33
```

5.3 Assignment

Kerf uses the colon (:) as an assignment operator, unlike the convention of “=” from many other programming languages. SQL uses = as a comparison operator, JSON uses : as an assignment operator in map literals and Kerf syntax attempts to be a superset of both JSON and SQL.

Values may be assigned to variables with : and retrieved by using the variable name. Note that uninitialized variables behave as containing an empty map:

```
Kerf> a
{}
Kerf> a: 3 7 19
[3, 7, 19]
Kerf> a
[3, 7, 19]
Kerf> a[1]
7
```

Assignment may be combined with indexing to assign to specific cells of a list or keys of a map:

```
Kerf> a: range 4
[0, 1, 2, 3]
Kerf> a[1]: 99
[0, 99, 2, 3]
Kerf> a
[0, 99, 2, 3]
Kerf> b: {bravo: 3, tango: 6};
Kerf> b["bravo"]: 99
{bravo:99, tango:6}
```

Note Kerf’s copy-on-write semantics:

```
Kerf> a: 0 1 2 3;
Kerf> b: a
[0, 1, 2, 3]
Kerf> b[1]:99
[0, 99, 2, 3]
Kerf> a
[0, 1, 2, 3]
```

As with indexing, it is possible to perform collective “spread” assignment:

```
Kerf> a: 8 take 0
[0, 0, 0, 0, 0, 0, 0, 0]
Kerf> a[1 2 5]:99
[0, 99, 99, 0, 0, 99, 0, 0]
Kerf> b: 8 take 0
[0, 0, 0, 0, 0, 0, 0, 0]
Kerf> b[1 2 5]:11 22 33
[0, 11, 22, 0, 0, 33, 0, 0]
```

It is also possible to perform compound assignment, treating `:` like a combinator which takes a binary function as a left argument and applies the old value and the right argument to this function before performing the assignment:

```
KeRF> a: 0 0 0
      [0, 0, 0]
KeRF> a#: 99
      [0, 0, 0, 99]
KeRF> a join: 47
      [0, 0, 0, 99, 47]
```

Compound assignment can be combined with indexing. Be warned, this can get confusing fairly quickly:

```
KeRF> a: 0 0 0
      [0, 0, 0]
KeRF> a[1] += 4
      [0, 4, 0]
KeRF> a[0 2] += 1
      [1, 4, 1]
KeRF> a[0 2] #: 55
      [[1, 55], 4, [1, 55]]
KeRF> a[0 2] #: 3 4
      [[1, 55, 3], 4, [1, 55, 4]]
```

To modify an element of a multidimensional structure, use multiple indexing expressions:

```
KeRF> a: [11 22, 33 44]
      [[11, 22],
       [33, 44]]
KeRF> a[0] [1] : 99
      [[11, 99],
       [33, 44]]
KeRF> a[0 1] [0] #: 0
      [[[11, 0], 99],
       [[33, 0], 44]]
```

5.4 Control Structures

Kerf has a familiar, simple set of general-purpose control structures. Parentheses and curly braces are *never* optional for control structures.

5.4.1 Conditionals

Kerf has a C-style if statement with optional `else if` and `else` clauses. Like the C ternary operator (`?:`), Kerf if statements can be used as part of an expression. Each curly-bracketed clause returns the value of its last expression.

```
Kerf> if (2 < 3) { 25 } else { 32 }  
25  
Kerf> if (2 > 3) { 25 } else { 32 }  
32
```

In Kerf, newlines are statement separators. Conditional statements spread across multiple lines must adhere to a specific, consistent indentation style:

```
if (a < b) {  
    c : 100  
} else if (a == b) {  
    c : 200  
} else {  
    c : 400  
}
```

Multiline conditionals *must not* be written with a newline before `else if` or `else` clauses. Implied (and undesirable) statement separators are shown in **red**:

```
if (a < b) {  
    c : 100  
};  
else if (a == b) {  
    c : 200  
};  
else {  
    c : 400  
}
```

Another incorrect indentation style:

```
if      (a < b) { c : 100 };  
else if (a == b) { c : 200 };  
else           { c : 400 };
```

5.4.2 Loops

Kerf provides a C-style for loop. The header consists of an initialization expression, a predicate and an updating expression. Note that the for loop itself returns null:

```
KerF> for (i: 0; i < 4; i: i+1) { display 2*i }  
0  
2  
4  
6  
  
KerF>
```

Kerf also provides a C-style while loop. Note how in this example the loop returns its final calculation:

```
KerF> t: 500; while(t > 32) { display t; t: floor t/2 }  
500  
250  
125  
62  
31  
  
KerF>
```

If you simply want to repeat an expression a fixed number of times, use do:

```
KerF> do (3) { display 42; 43 }  
42  
42  
42  
43  
  
KerF>
```

If it is necessary to prematurely exit a loop, break the loop into its own function and use return.

Combinators and built-in functions can be substituted for loops in many situations:

```
KerF> display mapdown range(0, 8, 2);  
0  
2  
4  
6  
  
KerF> {[t] t>32 } {[t] display t; floor t/2 } converge 500  
500  
250  
125  
62  
31  
  
KerF> display mapdown take(3, 42);  
42  
42  
42
```

5.4.3 Function Declarations

Functions can be declared using the `function` or `def` keywords and providing a parenthesized argument list. The final statement in a function body will be implicitly returned, and at any point in a function body you can instead use the `return` keyword to explicitly return.

```
function is_even(n) {  
    return (n % 2) == 0  
}  
  
def divisible(a, b) {  
    return (a % b) == 0  
}
```

It is also possible to define anonymous functions as part of an expression. Some languages refer to these as *lambdas*, in reference to the lambda calculus, a formal model of computation based on the manipulation of anonymous functions. Anonymous functions are enclosed in curly brackets and may provide a square-bracketed (`[` and `]`) argument list:

```
KeRF> {[a, b] 2*a+b }(3, 5)  
16
```

Storing a lambda in a variable is precisely equivalent to defining a function with `function` or `def`.

```
KeRF> divisible: {[a, b] (a % b) == 0 }  
{[a, b] (a % b) == 0}  
KeRF> divisible(6, 3)  
1  
KeRF> divisible(7, 2)  
0
```

KeRF uses *lexical scope*. This means that when variables are referenced, the *definition textually closest* to the reference will be used:

```
x: 35  
function outer_1() {  
    x: 25  
    function inner() {  
        return x  
    }  
    return inner;  
}  
  
function outer_2() {  
    x: 15  
    l: outer_1()  
    x: 45  
    return l()  
}  
display outer_2()
```

This example will print 25, because `inner` captures the definition of `x` in `outer_1` when it is created. The definitions of `x` in `outer_2` are not used when this function is evaluated, nor is the top-level definition of `x`.

6 SQL

Kerf understands SQL (Structured Queried Language), a popular programmatic interface for relational databases. You can blend SQL-style queries with imperative statements and access the full range of Kerf predicates and logical operators while filtering and selecting results.

SQL keywords are not case-sensitive, but for clarity the following examples will use uppercase exclusively. When describing the syntax of SQL statements, sections which can contain table names, field names or other types of subexpressions will be shown in **bold**. If a section is optional, it will be enclosed in bold square brackets (**[** and **]**). For example:

```
SELECT fields [AS name] FROM table ...
```

6.1 INSERT

INSERT is the simplest type of SQL statement. It is used for creating or appending to tables. INSERT can perform single or bulk insertions, the latter of which is much more efficient.

```
INSERT INTO table VALUES data
```

table can be a table literal or the name of a variable containing a table. In the latter case, the table will be modified in-place. **data** can be a list, matrix, map or table.

A single insertion can take values from a comma-separated, parenthesized list- this is special SQL syntax. Alternatively, use an ordinary square-bracketed Kerf list. You can also use a map, which more explicitly shows column names in the source data.

```
KerF> INSERT INTO {{name, email, level}} VALUES ("bob", "b@ob.com", 7)
```

name	email	level
bob	b@ob.com	7

```
KerF> INSERT INTO {{name, email, level}} VALUES ["bob", "b@ob.com", 7]
```

name	email	level
bob	b@ob.com	7

```
KerF> INSERT INTO {{name, email, level}} VALUES {name: "bob", level: 7, email: "b@ob.com"}
```

name	email	level
bob	b@ob.com	7

Note that a map which is to be inserted must have **exactly** the same key set as the destination table:

```
KeRF> INSERT INTO {{name, email, level}} VALUES {name: "bob", level: 7}

INSERT INTO {{name, email...
^
Length error
KeRF> INSERT INTO {{name, email, level}} VALUES {name: "bob", level: 7, hobbies: "needlepoint"}

INSERT INTO {{name, email...
^
Column error
```

Bulk insertions require either a matrix or a table:

```
KeRF> employees: [{"bob","alice","jerry"}
                  ["b@ob.com", "alice@gmail.com", "jerry@zombo.com"]
                  [7, 9, 43]];
KeRF> INSERT INTO {{name, email, level}} VALUES employees
```

name	email	level
bob	b@ob.com	7
alice	alice@gmail.com	9
jerry	jerry@zombo.com	43

An empty table will accept an INSERT from any map or table. List or matrix elements will be assigned default column names:

```
KeRF> INSERT INTO {{}} VALUES {legume: "Black Bean", dish: "Casserole"}
```

legume	dish
Black Bean	Casserole

```
KeRF> INSERT INTO {{}} VALUES employees
```

col	col1	col2
bob	b@ob.com	7
alice	alice@gmail.com	9
jerry	jerry@zombo.com	43

6.2 DELETE

DELETE is used for removing rows from a table. Kerf's columnar representation of tables means that a DELETE runs in linear time with respect to the number of rows in the table. Keep this in mind, and avoid repeated DELETES over large (on-disk) datasets.

```
DELETE FROM table [ WHERE condition ]
```

table can be a table literal or the name of a variable containing a table. In the latter case, the table will be modified in-place. **condition** can be any Kerf expression, using the names of columns from **table** as variables. For more information about WHERE, see the discussion of SELECT.

If provided a reference to a table stored in a variable, DELETE will return the name of that variable. Otherwise, it will return the modified table itself:

```
Kerf> t: {{a:range(5000), b:rand(5000, 100.0)}}
```

a	b
0	16.4771
1	27.3974
2	28.3558
3	12.2126
4	45.1148
5	81.5326
6	95.726
7	38.1769
.	..

```
Kerf> count t  
5000
```

```
Kerf> DELETE FROM t WHERE b between [0, 50]  
"t"
```

```
Kerf> count t  
2474
```

```
Kerf> DELETE FROM {{a:range(10), b:rand(10, 100.0)}} WHERE (a%2) = 1
```

a	b
0	75.3017
2	67.3
4	19.4571
6	66.3412
8	66.116

As you might expect, if you don't use a WHERE clause, DELETE will remove all the rows of a table:

```
Kerf> DELETE FROM t  
"t"
```

```
Kerf> t
```

a	b

6.3 SELECT

SELECT performs queries. It can be used to extract, aggregate or transform the contents of tables, producing new tables.

```
SELECT fields [ AS name ] FROM table
    [ WHERE condition ]
    [ GROUP BY aggregate ]
```

In its simplest form, SELECT can be used to slice a desired set of columns out of a table:

KeRF> `people`

name	age	gender	job
Hamilton Butters	37	M	Janitor
Emma Peel	29	F	Secret Agent
Jacques Maloney	48	M	Private Investigator
Renee Smithee	31	F	Programmer
Karen Milgram	16	F	Student
Chuck Manwich	29	M	Janitor
Steak Manhattan	18	M	Secret Agent
Tricia McMillen	29	F	Mathematician

KeRF> `SELECT name, gender FROM people`

name	gender
Hamilton Butters	M
Emma Peel	F
Jacques Maloney	M
Renee Smithee	F
Karen Milgram	F
Chuck Manwich	M
Steak Manhattan	M
Tricia McMillen	F

Selected columns can be renamed by specifying an AS clause for each:

KeRF> `SELECT age AS person_age, gender AS sex FROM people`

person_age	sex
37	M
29	F
48	M
31	F
16	F
29	M
18	M
29	F

The wildcard `*` can be used to refer to all columns in a table. It is also possible to use a variety of collective functions like `count` or `avg` when selecting columns. When calculated columns are not given a name explicitly via `AS`, a default name will be supplied.

```
KerF> SELECT count(*), sum(age), avg(age) AS average_age FROM people
```

col	age	average_age
8	237	29.625

It is possible to reference user-defined functions in a `SELECT`, but if they are not atomic you may need to explicitly apply them to elements of a column using combinators:

```
KerF> revname: {[x] p:explode(`" ",x); implode(" ", reverse p)};
KerF> from_dogyears: {[x] if (x < 21) { x/10.5 } else { x/4 }};
KerF> SELECT revname mapdown name, from_dogyears mapdown age AS dog_age FROM people
```

name	dog_age
Butters, Hamilton	9.25
Peel, Emma	7.25
Maloney, Jacques	12.0
Smithee, Renee	7.75
Milgram, Karen	1.52381
Manwich, Chuck	7.25
Manhattan, Steak	1.71429
McMillen, Tricia	7.25

6.3.1 WHERE

The `WHERE` clause permits filtering of results. Only rows which adhere to the constraints given as the **condition** will be returned:

```
KerF> SELECT * FROM people WHERE age > 30
```

name	age	gender	job
Hamilton Butters	37	M	Janitor
Jacques Maloney	48	M	Private Investigator
Renee Smithee	31	F	Programmer

You can form a *conjunction* with several conditions by separating them with commas. Given a conjunction, the result is only selected if all conditions are satisfied. This is equivalent to performing a logical *AND*:

```
KerF> SELECT * FROM people WHERE gender = "M", age > 30
```

name	age	gender	job
Hamilton Butters	37	M	Janitor
Jacques Maloney	48	M	Private Investigator

You can also form a conjunction by using the Kerf and operator, but in this case you *must* parenthesize subexpressions. Remember: Kerf evaluates expressions right to left unless otherwise parenthesized, so `a = b and c > d` is equivalent to `a = (b and (c > d))`:

```
KerF> SELECT * FROM people WHERE gender = "M" and age > 30
```

```
gender = "M" and age > 30
```

```
^
```

Type error

```
KerF> SELECT * FROM people WHERE (gender = "M") and (age > 30)
```

name	age	gender	job
Hamilton Butters	37	M	Janitor
Jacques Maloney	48	M	Private Investigator

6.3.2 GROUP BY

The GROUP BY clause can be used to gather together sets of rows which match on a particular column. In its simplest form, it behaves somewhat like the built-in function part. Note that the grouped-by column is included in the results:

```
KerF> SELECT name, age FROM people GROUP BY job
```

job	name	age
Janitor	[Hamilton Butters, Chuck Manwich]	[37, 29]
Secret Agent	[Emma Peel, Steak Manhattan]	[29, 18]
Private Investigator	[Jacques Maloney]	[48]
Programmer	[Renee Smithee]	[31]
Student	[Karen Milgram]	[16]
Mathematician	[Tricia McMillen]	[29]

You may often want to use collective functions to reduce each list of results:

```
KerF> SELECT count(name) AS num, avg(age) FROM people GROUP BY job
```

job	num	age
Janitor	2	33.0
Secret Agent	2	23.5
Private Investigator	1	48.0
Programmer	1	31.0
Student	1	16.0
Mathematician	1	29.0

6.4 UPDATE

UPDATE modifies a table in-place, altering the values of some or all columns of rows which match a query.

```
UPDATE table SET assignments
    [ WHERE condition ]
    [ GROUP BY aggregate ]
```

In its simplest form, UPDATE transforms or reassigns one or more of the columns of the table. The right side of each clause of **assignments** can be any Kerf expression. Note that the SET clause can use the symbols = or : to represent assignment, for compatibility with familiar SQL engines. In practice, favor using : to avoid ambiguity.

```
KeRF> UPDATE {{a:10 20 30}} SET a=a*2
```

a
20
40
60

```
KeRF> UPDATE {{a:10 20, b: 40 50}} SET a:11+a, b=5
```

a	b
21	5
31	5

```
KeRF> UPDATE {{a:0 1 2, b: 3 4 5}} SET b:5
```

a	b
0	5
1	5
2	5

Assignments are carried out left to right:

```
KeRF> UPDATE {{a:"First", b:"Second"}} SET a=b, b=a
```

a	b
Second	Second

An UPDATE can be restricted to only modify a specific subset of rows by using a WHERE clause, as in SELECT:

```
KeRF> UPDATE {{a:0 1 2, b:7 3 7}} SET b=99 WHERE b=7
```

a	b
0	99
1	3
2	99

6.5 Joins

Kerf provides built-in functions `left_join` and `asof_join` which can be used to align and combine tables:

```
KerF> livesin
```

name	nationality
Hamilton Butters	USA
Emma Peel	UK
Jacques Maloney	France
Renee Smithee	France
Karen Milgram	USA
Chuck Manwich	Canada
Tricia McMillen	UK

```
KerF> SELECT name, age, nationality FROM left_join(people, livesin, "name")
```

name	age	nationality
Hamilton Butters	37	USA
Emma Peel	29	UK
Jacques Maloney	48	France
Renee Smithee	31	France
Karen Milgram	16	USA
Chuck Manwich	29	Canada
Steak Manhattan	18	null
Tricia McMillen	29	UK

6.5.1 Left Join

A left join includes every row of the left table (x), and adds any additional columns from the right table (y) by matching on some key column (z). Added columns where there is no match on z will be filled with type-appropriate null values as generated by `type_null`.

```
KeRF> t: {{a:1 2 2 3, b:10 20 30 40}}
```

a	b
1	10
2	20
2	30
3	40

```
KeRF> u: {{a:2 3, c:1.5 3}}
```

a	c
2	1.5
3	3.0

```
KeRF> left_join(t, u, "a")
```

a	b	c
1	10	nan
2	20	1.5
2	30	1.5
3	40	3.0

If z is a list, require a match on several columns:

```
KeRF> u: {{a:2 3, b:30 40, c:1.5 3}};
```

```
KeRF> left_join(t, u, ["a","b"])
```

a	b	c
1	10	nan
2	20	nan
2	30	1.5
3	40	3.0

If `z` is a map, associate columns from `x` as keys with columns from `y` as values, permitting joins across tables whose column names differ.

```
KeRF> u: {{z:2 3, c:1.5 3}};  
KeRF> left_join(t, u, {'a':'z'})
```

a	b	c
1	10	nan
2	20	1.5
2	30	1.5
3	40	3.0

6.5.2 Asof Join

Behaves as `left_join` for the first three arguments. The fourth argument is a string, list or map indicating columns which will match if the values in `y` are less than or equal to `x`. Often this operation is applied to timestamp columns, but it works for any other comparable column type.

```
KeRF> t: {{a: 1 2 2 3, b: 10 20 30 40}}
```

a	b
1	10
2	20
2	30
3	40

```
KeRF> u: {{b: 19 17 32 8, c: ["A","B","C","D"]}}
```

b	c
19	A
17	B
32	C
8	D

```
KeRF> asof_join(t, u, [], "b")
```

a	b	c
1	10	D
2	20	A
2	30	A
3	40	C

6.6 Limiting

If you wish to retrieve the first n items of a query, as in a SQL LIMIT clause, you can use first:

```
KeRF> first(2, SELECT name, age FROM people WHERE gender = "F")
```

name	age
Emma Peel	29
Renee Smithee	31

But beware- if the result has fewer than n rows, this approach will replicate them:

```
KeRF> first(2, SELECT name, age FROM people WHERE name = "Emma Peel")
```

name	age
Emma Peel	29
Emma Peel	29

A better approach is to define a new function which takes the minimum of n and the length of the result:

```
KeRF> limit: {[n, t] first(min(count t, n), t)};  
KeRF> limit(2, SELECT name, age FROM people WHERE name = "Emma Peel")
```

name	age
Emma Peel	29

6.7 Ordering

If you wish to sort tables along a column, as in a SQL ORDER BY clause, you can use the built-in functions ascend or descend along with indexing:

```
KeRF> people[ascend SELECT job FROM people]
```

name	age	gender	job
Hamilton Butters	37	M	Janitor
Chuck Manwich	29	M	Janitor
Tricia McMillen	29	F	Mathematician
Jacques Maloney	48	M	Private Investigator
Renee Smithee	31	F	Programmer
Emma Peel	29	F	Secret Agent
Steak Manhattan	18	M	Secret Agent
Karen Milgram	16	F	Student

6.8 Performance

WHERE clauses have a special understanding of certain Kerf verbs and can achieve significant performance boosts in the right circumstances.

```
KerF> n: 200000;
KerF> i: range(n);
KerF> v: rand(n, 100.0);
KerF> iv: indexed v;
KerF> find: {[x] SELECT count(*) FROM {{i:i, v:x}} WHERE v < 23.7};

KerF> timing 1;
KerF> find v;
      15 ms
KerF> find iv;
      6 ms
```

For best results, order WHERE conjunctions to perform the largest reduction of data first, or take advantage of indexed or enum columns as early as possible:

```
KerF> n: 200000;
KerF> t: {{a: range(n), b: rand(n, 100.0), c: rand(n, 6)}}}
```

a	b	c
0	82.268	2
1	80.5227	0
2	13.5797	1
3	80.2291	3
4	61.5329	3
5	67.2546	1
6	64.5684	5
7	29.7027	2
.	..	.

```
KerF> timing 1;
KerF> SELECT * FROM t WHERE b > 50, c = 1;
      25 ms
KerF> SELECT * FROM t WHERE c = 1, b > 50;
      13 ms
KerF> SELECT * FROM t WHERE (c = 1) and (b > 50);
      14 ms
```

7 Input/Output

Kerf provides a rich set of built-in IO functions for displaying, serializing, importing and exporting data. Beyond the capabilities described here, Kerf provides a general purpose foreign-function interface- see **FFI**.

7.1 General I/O

The function `out(x)` will print a string `x` to standard output. Non-string values are ignored:

```
KerF> out "foo"
foo
KerF> out 65
KerF>
```

The function `display(x)` will print a display representation of data to standard output. A key difference between calling this function from the REPL and using the REPL's natural value printing is that `display` will print the entire result:

```
KerF> range 50
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19,
20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37,
38, 39, 40, ...]
KerF> display range 50
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19,
20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37,
38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49]
```

The function `shell(x)` will execute a string `x` containing a shell command as if from `/bin/sh -c x` and return the lines of the result:

```
KerF> out implode("\n", shell("cal"))
November 2015
Su Mo Tu We Th Fr Sa
1 2 3 4 5 6 7
8 9 10 11 12 13 14
15 16 17 18 19 20 21
22 23 24 25 26 27 28
29 30
KerF> shell("echo Hello, World!")
["Hello, World!"]
```

7.2 File I/O

The function `dir_ls` lists the files and directories at a path. It can optionally provide full path names:

```
KeRF> dir_ls("/Users/john/Sites")
[".DS_Store", ".localized", "images", "index.html", "subforum.php"]
KeRF> dir_ls("/Users/john/Sites", 1)
["\\Users\\john\\Sites\\.DS_Store"
 "\\Users\\john\\Sites\\.localized"
 "\\Users\\john\\Sites\\images"
 "\\Users\\john\\Sites\\index.html"
 "\\Users\\john\\Sites\\subforum.php"]
```

The function `lines(filename, n)` loads lines from a plain text file into a list of strings. The argument `n` is optional, and specifies the maximum number of lines to read.

```
KeRF> lines("example.txt")
["First line", "Second line", "Third line"]
KeRF> lines("example.txt", 2)
["First line", "Second line"]
```

The function `write_text(filename, x)` writes a raw string to a file, returning the number of bytes written. If `x` is not already a string it will be converted to one as by `json_from_kerf`. To perform the inverse of `lines()`, use `write_text(filename, implode("\n", x))`.

```
KeRF> write_text("example.txt", 5)
1
KeRF> write_text("example.txt", 99)
2
KeRF> shell("cat example.txt")
["99"]
```

Kerf has a proprietary binary serialization format. The function `write_to_path(filename)` writes a Kerf object to a file, creating it if necessary, and returns 0 if the operation was successful. Objects can then be reconstituted from binary files by calling `read_from_path(filename)`.

```
KeRF> write_to_path("example.bin", 23 24 25)
0
KeRF> shell("wc -c example.bin")
["    64 example.bin"]
KeRF> out implode("\n", shell("hexdump example.bin"))
00000000 06 90 00 fe 01 00 00 00 03 00 00 00 00 00 00 00
00000100 17 00 00 00 00 00 00 00 18 00 00 00 00 00 00 00
00000200 19 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000300 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000400
KeRF> read_from_path("example.bin")
[23, 24, 25]
```

Kerf also has built-in functions which make it easy to read and write common tabular data formats. `read_table_from_delimited_file` loads data, `write_delimited_file_from_table` writes data, and several wrapper functions are available which read and write CSV or TSV files using this functionality.

Kerf can load data from text files with fixed-width columns with `read_table_from_fixed_file`. See the built-in function reference discussion of this function for additional details.

7.3 Striped Files

Kerf is also capable of working with “striped” files. This serialization format writes each nested component of an object (the items of a list, or the columns of a table, for example) as a separate file within a directory structure. Striped files are more space-efficient for sparse structures, and permit significantly faster appends.

To create a striped file from an object, use `write_striped_to_path(path, x)`:

```
Kerf> write_striped_to_path("stripe_example", {{a: 1 2 3, b: 4 5 6}});
Kerf> out "\n" implode shell "tree stripe_example"
stripe_example
├── 2
│   ├── 0
│   │   └── base.dat
│   ├── 1
│   │   └── base.dat
│   └── base.dat
└── base.dat

3 directories, 4 files
```

To load a striped file, use `read_striped_from_path(path)`:

```
Kerf> read_striped_from_path("stripe_example")
```

a	b
1	4
2	5
3	6

For high-performance data storage, Kerf provides a set of alternatives to the “`read.table_`”-prefixed methods which begin with “`build.table_`”, like `build.table_from_csv`. Instead of simply reading a file into memory, these routines load data into a memory-mapped, disk-backed data structure, appending to any pre-existing file. These routines will use striped files preferentially, but also support conventional serialized Kerf objects. The `build.table_` functions allow you to work efficiently with large data files which may not fit in memory at once.

7.4 Network I/O (IPC)

Kerf has a specialized remote procedure call system built on TCP which permits distributing tasks across multiple processes or machines. This is sometimes described as Inter-Process Communication (IPC). To spawn a Kerf process which listens for IPC calls, invoke `kerf` with the `-p` command-line argument and specify a TCP port number:

```
> ./kerf -p 1234
Kerf Server listening on port: 1234
KeRF>
```

The function `open_socket(host, port)` opens a connection to a remote Kerf instance at hostname `host` and listening on port and returns a connection handle. Both `host` and `port` must be strings. Once opened, a connection handle may be closed via `close_socket(handle)`.

```
KeRF> open_socket("localhost", "1234")
4
```

The function `send_async(handle, x)` will send a string `x` to a remote Kerf instance without waiting for a reply. `x` will be eval'd on the remote server, returning 1 on a successful send.

The function `send_sync(handle, x)` will send a string `y` to a remote Kerf instance, waiting for a reply. `y` will be eval'd on the remote server, and the result will be returned.

```
KeRF> c: open_socket("localhost", "1234")
4
KeRF> send_async(c, "foo: 2+3")
1
KeRF> foo
{}
KeRF> send_sync(c, "[foo, foo]")
[5, 5]
KeRF> send_sync(c, "sum range 1000")
499500
```

During IPC execution, the constant `.Net.client` contains the current client's unique handle:

```
KeRF> open_socket("localhost", "10101")
6
KeRF> send_sync(6, ".Net.client")
6
KeRF> .Net.client
0
```

If defined, an IPC server will call the single-argument function `.Net.on_close` with a client handle when that client closes its connection:

```
KeRF> .Net.on_close: {[x] out 'client closed: ' join (string x) join '\n'};
KeRF>
server: new connection from 127.0.0.1 on socket 6
client closed: 6
```

8 Foreign Function Interface

Kerf provides a Foreign Function Interface (FFI) which makes it possible to call out to C libraries from Kerf and call into Kerf from C. The FFI permits users to supplement Kerf with new IO capabilities, make use of pre-existing libraries and fine-tune high performance applications. In this section we will discuss how to use this facility and describe the data structures and the native API exposed by the Kerf executable.

8.1 Overview

To compile a dynamic library for use with Kerf, you will need `kerf_api.h`, which should be included with the Kerf binary and other documentation. This contains the signatures of exposed `kerf_api_` functions you can use in building dynamic libraries as well as a number of useful constants. If your C source code is stored in `example.c` and this header is in the current directory, compile it from the command line as follows:

For OSX:

```
>cc -m64 -flat_namespace -undefined suppress -dynamiclib example.c -o example.dylib
```

The `-m64` option requests code be generated for a 64-bit architecture. The `-flat_namespace` and `-undefined suppress` are necessary to create a dynamic library which can be linked against the Kerf binary at runtime—this is how your dynamic library can access built-in utility methods from the interpreter itself. The `-o` option specifies a name for the output file, which we suffix with `.dylib` by convention.

For Linux:

```
>cc -m64 -shared -fpic example.c -o example.dylib
```

Linux also requires `-m64`, but by contrast uses `-fpic` and `-shared` to build a shared library which can be dynamically linked to Kerf at runtime. While Linux generally uses `.so` rather than `.dylib` suffixes for shared objects, we will stick to the OSX naming convention for these examples for the sake of consistency.

With this dynamic library compiled, the `dlopen` built-in function can be used to load the library and create a wrapper function for a specific function from the library which can then be invoked in Kerf as if it were an ordinary Kerf function. The Kerf FFI does not marshal Kerf's native data representations into C types—if another representation is desired you must close the gap in your code. This approach simplifies the FFI implementation and permits very efficient data exchange, but requires a great deal of care and attention from library authors.

All values in Kerf are internally stored in KERF structures. This struct is what is called a *tagged union*- some fields have different interpretations dependent on the value of a *tag field*. In our case, the typecode field `t` permits this discrimination.

```
typedef struct kerf0 {
    char m;      // log_2 of memory slab size (internal use only)
    char a;      // attribute flags
    char h;      // log_2 of header subslice size (internal use only)
    char t;      // typecode, as in kerf_type
    int32_t r;   // reference count

    union {
        int64_t i; // value of a Kerf integer or stamp
        double f; // value of a Kerf float
        char c; // value of a Kerf char
        char* s;
        struct kerf0* k;

        struct {
            int64_t n; // in a vector or list, the number of elements
            char g[]; // in a vector or list, the elements themselves
        };
    };
} *KERF, KERFO;
```

Some of the fields of a KERF structure are for internal use by the interpreter and should not be manipulated directly in well-behaved dynamic libraries. The next section will describe the internal representation of various Kerf types in terms of KERF structures. Three fields are of general interest:

- The `r` field contains a *reference count* for the object. Examining this field explicitly is sometimes useful for debugging memory corruption problems. Your code should treat it as read-only.
- The `t` field indicates the object's *typecode*, identical to the typecodes given by `kerf_type`. Negative values indicate vector types.
- The `a` field contains *attribute flag* bits which provide the interpreter with metadata that alters the behavior of values or permits runtime optimizations.

Functions called from Kerf will always take some number of KERF structures as arguments and return a KERF structure. In the event that your function doesn't produce a meaningful return value, return 0- Kerf will coerce this into the equivalent of null.

8.2 Reference Counting

Kerf manages dynamically allocated memory using a strategy called *reference counting*. Each KERF object contains a counter of how many references to the object exist throughout the runtime. Whenever an object is used or stored, its reference count is incremented. When an object is no longer needed (perhaps a function using the object completes and returns) this count is decremented. When an object has a reference count of 0, the Kerf runtime knows that it is free to deallocate the object and reclaim the associated memory.

Reference counting is contrasted with *garbage collection*, another popular method of managing dynamic memory. A garbage collector actively *searches* for references to objects and automatically frees objects when they are no longer reachable. Garbage collectors remove the need to explicitly maintain reference counts, but generally incur runtime overhead for performing their heap scanning. Kerf uses reference counting because it has low runtime overhead and, by giving programmers more explicit awareness of memory management, permits more predictable runtime performance.

If the reference counts on an object are not maintained properly, a variety of bugs can occur. If references are not *released* when an object is no longer needed, memory will *leak*- it is no longer useful but cannot be reclaimed for other purposes. In long-running processes this can eventually exhaust system memory. If references are not *retained* when an object is still needed, they can be deallocated out from under your code, allowing data to be overwritten and creating a variety of potential intermittent failures.

Some objects require special consideration. For efficiency purposes, Kerf will sometimes allocate objects which can be passed as arguments to your dynamic libraries on the *stack* instead of its reference-counted *heap*. Stack-allocated objects typically have a reference count of -1 and cannot be retained indefinitely without copying. These and other special cases mean that simply decrementing or incrementing the *r* field of a KERF object is not sufficient for controlling reference counts correctly. The Kerf native API exposes two functions- `kerf_api_release` and `kerf_api_retain`- for this purpose which perform all the necessary bookkeeping. Do *not* directly modify the *r* field of KERF objects!

8.2.1 Releasing

As a general rule, if a KERF object is allocated in a function (as with any of the `kerf_api_new_` functions) it must either be returned from that function or released with `kerf_api_release` before the function returns. The following routine leaks a `KERF_CHARVEC`. Can you see why?

```
KERF leaky_demo(KERF monad) {  
    return kerf_api_call_monad(monad, kerf_api_new_charvec("Hello!"));  
}
```

A corrected version would be as follows:

```
KERF leaky_demo(KERF monad) {  
    KERF string = kerf_api_new_charvec("Hello!");  
    KERF result = kerf_api_call_monad(monad, string);  
    kerf_api_release(string);  
    return result;  
}
```

Any objects which were allocated inside `kerf_api_call_monad` are the responsibility of that routine, and the result it returns can be freely returned from our routine. Objects passed *into* our routine are the responsibility of the code which called our routine, and will be automatically deallocated unless we explicitly retain them. This idea of *local responsibility* is the key to understanding reference counting.

8.2.2 Retaining

In many situations, simply releasing objects you allocate is sufficient, but in some situations we need to explicitly retain an object to ensure it is not deallocated too soon. Consider this pass-through function which prints information about reference counts:

```
KERF show_refcount(KERF x) {  
    printf("incoming reference count: %d\n", x->r);  
    KERF result = kerf_api_retain(x);  
    printf("outgoing reference count: %d\n", result->r);  
    return result;  
}
```

The calling function which passed the arguments into our dynamic library function is responsible for discarding those arguments when our function returns. If we want to return a value that was originally passed to us as an argument, we must retain it to counteract this.

```
KerF> refcount: dlopen("example.dylib", "show_refcount", 1);  
KerF> refcount(5)  
incoming reference count: -1  
outgoing reference count: 1  
5  
KerF> refcount([1,2,3])  
incoming reference count: 2  
outgoing reference count: 3  
[1, 2, 3]  
KerF> a: 4 5 6;  
KerF> refcount(a)  
incoming reference count: 3  
outgoing reference count: 4  
[4, 5, 6]
```

In the first example above, the value was stack-allocated, and the call to `kerf_api_retain` copied this value to the heap, returning a pointer to the new object. **Always** use the result of a call to `kerf_api_retain` instead of the original value.

Retaining is also important if you want to stash a value for later use. Here's an example of some dynamic library functions which hang onto state:

```
KERF stashed_value = 0;

KERF state_init(KERF value) {
    stashed_value = kerf_api_retain(value);
    return 0;
}
KERF state_double() {
    KERF result = stashed_value;
    stashed_value = kerf_api_new_int((result->i) * 2);
    return result;
}
KERF state_free(KERF value) {
    kerf_api_release(stashed_value);
    stashed_value = 0;
    return 0;
}
```

We retain the initial value passed in, since it would otherwise be destroyed when the function returns, as in the previous example. In subsequent calls to `state_double`, we can return the stashed value, relinquishing control of it. The new value we construct each time has a reference count of 1 and can also be returned when the time comes. Freeing our stashed state is included for completeness.

```
KerF> init: dlopen("example.dylib", "state_init", 1);
KerF> step: dlopen("example.dylib", "state_double", 0);
KerF> free: dlopen("example.dylib", "state_free", 0);
KerF> init(5);
KerF> step()
5
KerF> step()
10
KerF> step()
20
KerF> free();
```

8.3 Internal Representations

To work with KERF structures, it is necessary to understand certain details of how they are represented. The implementation of some data structures, however, is kept intentionally opaque. Even dynamic libraries should not depend on distinguishing the internal structure of an enumeration or an index from a mixed-type list, for example. For low-level operations, KERF structures can be directly consulted, and for operations on higher-level datatypes the native API exposes helper functions.

8.3.1 Integers

Kerf Integers have a typecode of `KERF_INT`, and store their values in the `i` field of a KERF structure as a signed 64-bit integer.

8.3.2 Floats

Kerf Floats have a typecode of `KERF_FLOAT`, and store their values in the `f` field of a KERF structure as an IEEE-754 double-precision floating point number.

8.3.3 Characters

Kerf Characters have a typecode of `KERF_CHAR`, and store their values in the `c` field of a KERF structure as UTF-8 characters.

8.3.4 Stamps

Kerf Timestamps have a typecode of `KERF_STAMP`, and store their values in the `i` field of a KERF structure as a signed 64-bit count of nanoseconds since Unix Epoch.

8.3.5 Vectors

Vectors store their data in the `g` field of a KERF structure as a raw, packed array, and use the `n` field to indicate the number of elements they contain. Vectors are typically slab-allocated, and may contain extra allocated space beyond the item count given in `n`, but do not assume this will be the case. See `kerf_api_new_kerf` for an example of how to create a vector from scratch.

8.4 Attribute Flags

The `a` field of a KERF structure is an integer which represents a vector of single-bit *flag* values. Each flag attached to a value represents a boolean piece of metadata which describes properties or configuration settings of the value. Kerf uses many attribute flags for internal bookkeeping, but a few of these flags represent information that could be generally useful for your dynamic libraries.

To test for the presence of a flag, bitwise AND the flag constant with the `a` field:

```
int flag_is_set = ((x->a) & KERF_ATTR_Y) != 0;
```

To set a flag, bitwise OR the flag constant into the `a` field:

```
(x->a) |= KERF_ATTR_Y;
```

To clear a flag, bitwise AND the `a` field with the bitwise complement of the flag constant. Think of it as “keep everything *except* this flag”:

```
(x->a) &= ~KERF_ATTR_Y;
```

The `KERF_ATTR_SORTED` flag indicates that the elements of a vector or list is sorted in ascending order. The same information can be queried in Kerf using the `sort_debug` built-in function.

```
KERF is_sorted(KERF list) {  
    return kerf_api_new_int((list->a & KERF_ATTR_SORTED) != 0);  
}
```

```
KerF> f: dload("example.dylib", "is_sorted", 1);  
KerF> f(1 2 3)  
1  
KerF> f(3 1 2)  
0
```

The `KERF_ATTR_BYTES` flag applies specifically to character vectors. If set, their contents will be pretty-printed as a series of hexadecimal bytes with a `0x` prefix instead of as characters.

```
KERF demo_attr_bytes_set(KERF charvec) {  
    charvec = kerf_api_retain(charvec); // returning a modified argument, must retain  
    charvec->a |= KERF_ATTR_BYTES;  
    return charvec;  
}
```

```
KerF> dload("example.dylib", "is_sorted", 1)("ABCDE")  
0x4142434445
```

8.5 Native API

The signatures of the following methods are described in a machine-readable form in `kerf_api.h`.

8.5.1 `kerf_api_append` - Append to List

KERF `kerf_api_append(KERF x, KERF y)`

Append an object `y` to the list `x`. If the original list can be referenced from elsewhere (if it is stored in a global, for example), you must retain it before appending. It is safe to call `kerf_api_append` on a list you construct yourself without any special work.

```
KERF append_demo(KERF base) {
    KERF tail = kerf_api_new_int(54);
    KERF list = kerf_api_append(kerf_api_retain(base), tail);
    kerf_api_release(tail);
    return list;
}
```

```
KerF> f: dlopen("example.dylib", "append_demo", 1);
KerF> f([])
[54]
KerF> f(1 2 3)
[1, 2, 3, 54]
KerF> f(5)
[5, 54]
KerF> a: 4 5 6;
KerF> f(a)
[4, 5, 6, 54]
KerF> a
[4, 5, 6]
```

Note that the value stored in the global `a` *was not modified*- we altered a copy.

8.5.2 `kerf_api_call_dyad` - Call Dyad

`kerf_api_call_dyad(KERF func, KERF x, KERF y)`

Call a dyadic (binary) function and return the result.

```
KERF dyad_demo(KERF dyad) {
    KERF a1 = kerf_api_new_int(3);
    KERF a2 = kerf_api_new_int(5);
    KERF ret = kerf_api_call_dyad(dyad, a1, a2);
    kerf_api_release(a1);
    kerf_api_release(a2);
    return ret;
}
```

```
KerF> dlopen("example.dylib", "dyad_demo", 1)({[x,y] [x,y,x+y,x*y]})
[3, 5, 8, 15]
```

8.5.3 kerf_api_call_monad - Call Monad

kerf_api_call_monad(KERF func, KERF x)

Call a monadic (unary) function and return the result.

```
KERF monad_demo(KERF monad) {
    KERF arg = kerf_api_new_int(17);
    KERF ret = kerf_api_call_monad(monad, arg);
    kerf_api_release(arg);
    return ret;
}
```

```
KerF> dload("example.dylib", "monad_demo", 1)({[x] [x,0,2*x,x*x]})
[17, 0, 34, 289]
```

8.5.4 kerf_api_call_nilad - Call Nilad

kerf_api_call_nilad(KERF func)

Call a niladic function (one which takes no arguments) and return the result.

```
KERF nilad_demo(KERF nilad) {
    KERF temp = kerf_api_call_nilad(nilad);
    kerf_api_show(temp);
    kerf_api_release(temp);
    return kerf_api_call_nilad(nilad);
}
```

```
KerF> a:[25]
[25]
KerF> dload("example.dylib", "nilad_demo", 1)({[] a[0]++; [3,2,a[0]]})
[3, 2, 26]    // printed by kerf_api_show
[3, 2, 27]    // returned
```

8.5.5 kerf_api_copy_on_write - Copy On Write

KERF kerf_api_copy_on_write(KERF x)

If x is referenced by multiple owners or stack-allocated, make a shallow copy and return this copy. Otherwise, it is safe to modify x in place. kerf_api_copy_on_write may free the object passed in, so if you don't want the original destroyed it must be retained first.

8.5.6 kerf_api_get - Kerf Get

KERF kerf_api_get(KERF x, KERF index)

Index into the KERF list x. Equivalent to the normal Kerf expression x[index].

```
KERF get_demo(KERF list, KERF index) {  
    KERF zero = kerf_api_new_int(0);  
    KERF first = kerf_api_get(list, zero);  
    kerf_api_show(first);  
    kerf_api_release(zero);  
    kerf_api_release(first);  
    return kerf_api_get(list, index);  
}
```

```
KerF> f: dload("example.dylib", "get_demo", 2);  
KerF> f("ABCD", 2)  
`"A"           // printed by kerf_api_show  
  `"C"           // returned  
KerF> f(1 3 7 10, 2 0 2)  
1  
[7, 1, 7]      // returned
```

8.5.7 kerf_api_interpret - Interpret String

KERF kerf_api_interpret(KERF charvec)

Execute a KERF_CHARVEC (string) as if via the built-in function eval.

```
KERF interpret_demo() {  
    KERF s1 = kerf_api_new_charvec("a: 1+2");  
    KERF s2 = kerf_api_new_charvec("2*range a");  
    KERF r1 = kerf_api_interpret(s1);  
    KERF r2 = kerf_api_interpret(s2);  
    kerf_api_release(s1);  
    kerf_api_release(s2);  
    kerf_api_show(r1);  
    kerf_api_show(r2);  
    kerf_api_release(r1);  
    kerf_api_release(r2);  
    return 0;  
}
```

```
KerF> dload("example.dylib", "interpret_demo", 0)()  
3  
[0, 2, 4]
```

8.5.8 kerf_api_len - Kerf Length

```
int64_t kerf_api_len(KERF x)
```

Determine the length of a KERF object `x`, as given by the built-in function `len()`. This routine is more general than consulting the `n` field of a KERF structure, and will produce sensible results for non-vector types.

```
KERF len_demo() {
    KERF string = kerf_api_new_charvec("Hello, C!");
    KERF number = kerf_api_new_int(43);
    KERF list   = kerf_api_new_list();
    printf("%d %d %d\n",
           (int)kerf_api_len(string),
           (int)kerf_api_len(number),
           (int)kerf_api_len(list)
    );
    kerf_api_release(string);
    kerf_api_release(number);
    kerf_api_release(list);
    return 0;
}
```

```
KerF> dlopen("example.dylib", "len_demo", 0)();
9 1 0
```

8.5.9 kerf_api_new_charvec - New Kerf String

```
KERF kerf_api_new_charvec(char* cstring)
```

A wrapper for `kerf_api_new_kerf()` which allocates and initializes a `KERF_CHARVEC` (string) from a null-terminated C string. The supplied string will be copied into the new vector; subsequent mutations will not propagate from one to the other.

8.5.10 kerf_api_new_float - New Kerf Float

```
KERF kerf_api_new_float(double n)
```

A wrapper for `kerf_api_new_kerf()` which allocates and initializes a single Kerf Float from an IEEE-754 double-precision floating point number `n`.

8.5.11 kerf_api_new_int - New Kerf Integer

```
KERF kerf_api_new_int(int64_t n)
```

A wrapper for `kerf_api_new_kerf()` which allocates and initializes a single Kerf Integer from a signed 64-bit C integer `n`.

8.5.12 `kerf_api_new_kerf` - New Kerf Object

`KERF kerf_api_new_kerf(char type, int64_t length)`

Allocate a raw KERF structure with a given type. For scalar types like `KERF_INT`, length should be 0. For vector types, length will be the number of entries the object will contain. This routine allocates space in Kerf's internal memory pool and initializes reference counts appropriately.

```
KERF intvec_demo(KERF base, KERF count) {
    KERF ret = kerf_api_new_kerf(KERF_INTVEC, count->i);
    for(int x = 0; x < count->i; x++) {
        ((int64_t*)(ret->g))[x] = (base->i) + x;
    }
    return ret;
}
```

```
KeRF> f: dlopen("example.dylib", "intvec_demo", 2);
KeRF> v: f(70000, 5)
      [70000, 70001, 70002, 70003, 70004]
KeRF> kerf_type_name v
      "integer vector"
```

8.5.13 `kerf_api_new_list` - New Kerf List

`KERF kerf_api_new_list()`

Allocate an empty list. Useful in combination with `kerf_api_append`.

```
KERF list_demo() {
    return kerf_api_new_list();
}
```

```
KeRF> v: dlopen("example.dylib", "list_demo", 0)()
      []
KeRF> kerf_type v
      6
KeRF> kerf_type_name v
      "list"
KeRF> len v
      0
```

8.5.14 kerf_api_new_map - New Kerf Map

KERF kerf_api_new_map()

Allocate an empty map.

```
KERF map_demo() {  
    return kerf_api_new_map();  
}
```

```
KeRF> v: dload("example.dylib", "map_demo", 0)()  
  
KeRF> kerf_type v  
7  
KeRF> kerf_type_name v  
"map"  
KeRF> len v  
1
```

8.5.15 kerf_api_new_stamp - New Kerf Timestamp

KERF kerf_api_new_stamp(int64_t nanoseconds)

A wrapper for `kerf_api_new_kerf()` which allocates and initializes a single Kerf Stamp from a signed 64-bit count of nanoseconds since Unix Epoch `n`.

```
KERF stamp_demo(KERF ns) {  
    KERF epoch = kerf_api_new_stamp(0);  
    kerf_api_show(epoch);  
    kerf_api_release(epoch);  
    return kerf_api_new_stamp(ns->i);  
}
```

```
KeRF> t: dload("example.dylib", "stamp_demo", 1)(3000000)  
00:00:00.000    // printed by kerf_api_show  
00:00:00.003    // returned  
KeRF> t[["year", "month", "day"]]  
[1970, 1, 1]
```

8.5.16 kerf_api_nil - Kerf Nil

KERF kerf_api_nil()

Allocate an object representing nil/null.

```
KERF nil_demo() {  
    return kerf_api_nil();  
}
```

```
KerF> v: dload("example.dylib", "nil_demo", 0)();  
  
KerF> kerf_type v  
5  
KerF> kerf_type_name v  
"null"  
KerF> len v  
1
```

8.5.17 kerf_api_release - Release Kerf Reference

void kerf_api_release(KERF x)

Reduce the reference count for a KERF object. Any objects which are allocated in a dynamic library call and not returned or otherwise stored must be manually released, or they will not be reclaimed.

8.5.18 kerf_api_retain - Retain Kerf Reference

KERF kerf_api_retain(KERF x)

Increase the reference count for a KERF object. This process may require making a copy of the source object.

8.5.19 kerf_api_set - Kerf Set

KERF kerf_api_set(KERF x, KERF index, KERF replacement)

Index into the KERF list x and replace the element with replacement, returning the modified list. Similar to the normal Kerf expression `x[index]:replacement`, but incapable of spread assignment. This operation will modify x in place.

```
KERF set_demo(KERF list, KERF index, KERF rep) {  
    return kerf_api_set(kerf_api_retain(list), index, rep);  
}
```

```
KerF> f: dload("example.dylib", "set_demo", 3);  
KerF> f(1 2 3, 1, 99)  
[1, 99, 3]
```

8.5.20 kerf_api_show - Show Kerf Object

KERF kerf_api_show(KERF x)

Print a prettyprinted representation of a KERF structure x to stdout, as displayed by the REPL, and return the structure unchanged. This routine primarily exists for debugging purposes.

```
KERF show_demo(KERF argument) {  
    kerf_api_show(argument);  
    KERF string = kerf_api_new_charvec("Hello, C!");  
    kerf_api_show(string);  
    kerf_api_release(string);  
    return 0;  
}
```

```
KeRF> dload("example.dylib", "show_demo", 1)(1 2 3);  
[1, 2, 3]  
"Hello, C!"
```

9 Built-In Function Reference

9.1 abs - Absolute Value

`abs(x)`

Calculate the absolute value of x. Atomic.

```
KeRF> abs -4 7 -2.19 NaN
[4, 7, 2.19, nan]
```

9.2 acos - Arc Cosine

`acos(x)`

Calculate the arc cosine (inverse cosine) of x, expressed in radians, within the interval [-1,1]. Atomic. The results of `acos` will always be floating point values.

```
KeRF> acos 0.5 -0.2 1
[1.0472, 1.77215, 0]
KeRF> cos(acos 0.5 -0.2 1 4)
[0.5, -0.2, 1, nan]
```

9.3 add - Add

`add(x, y)`

Calculate the sum of x and y. Fully atomic.

```
KeRF> add(3, 5)
8
KeRF> add(3, 9 15 -7)
[12, 18, -4]
KeRF> add(9 15 -7, 3)
[12, 18, -4]
KeRF> add(9 15 -7, 1 3 5)
[10, 18, -2]
```

The symbol `+` is equivalent to `add` when used as a binary operator:

```
KeRF> 2 4 3+9
[11, 13, 12]
```

9.4 and - Logical AND

`and(x, y)`

Calculate the logical *AND* of x and y. This operation is equivalent to the primitive function `min`. Fully atomic.

```
KeRF> and(1 1 0 0, 1 0 1 0)
[1, 0, 0, 0]
KeRF> and(1 2 3 4, 0 -4 9 0)
[0, -4, 3, 0]
```

The symbol & is equivalent to and when used as a binary operator:

```
KeRF> 1 1 0 0 & 1 0 1 0
[1, 0, 0, 0]
```

9.5 ascend - Ascending Indices

`ascend(x)`

For a list `x`, generate a list of indices into `x` in ascending order of the values of `x`.

```
KeRF> t:5 2 3 1
[5, 2, 3, 1]
KeRF> ascend t
[3, 1, 2, 0]
KeRF> t[ascend t]
[1, 2, 3, 5]
```

Strings are sorted in lexicographic order:

```
KeRF> ascend ["Orange","Apple","Pear","Aardvark","A"]
[4, 3, 1, 0, 2]
```

When applied to a map, `ascend` will sort the keys by their values and produce a list:

```
KeRF> ascend {"A":2, "B":9, "C":0}
["C", "A", "B"]
```

The symbol < is equivalent to `ascend` when used as a unary operator:

```
KeRF> <5 2 3 1
[3, 1, 2, 0]
```

9.6 asin - Arc Sine

`asin(x)`

Calculate the arc sine (inverse sine) of `x`, expressed in radians, within the interval `[-1,1]`. Atomic. The results of `asin` will always be floating point values.

```
KeRF> asin 0.5 -0.2 1
[0.523599, -0.201358, 1.5708]
KeRF> sin(asin 0.5 -0.2 1 4)
[0.5, -0.2, 1, nan]
```

9.7 asof_join - Asof Join

`asof_join(x, y, k1, k2)`

Perform a "fuzzy" left join. See **Joins**.

9.8 atan - Arc Tangent

atan(x)

Calculate the arc tangent (inverse tangent) of x, expressed in radians. Atomic. The results of atan will always be floating point values.

```
KeRF> atan 0.5 -0.2 1 4
[0.463648, -0.197396, 0.785398, 1.32582]
KeRF> tan(atan 0.5 -0.2 1 4)
[0.5, -0.2, 1, 4]
```

9.9 atlas - Atlas Of

atlas(map)

Create an atlas from a map. An atlas is the schemaless NoSQL equivalent of a table. Atlases are automatically indexed in such a way that all key-queries are indexed.

```
KeRF> atlas({name:["bob", "alice", "oscar"], id:[123, 421, 233]})
atlas[{name:["bob", "alice", "oscar"], id:[123, 421, 233]}]
```

9.10 atom - Is Atom?

atom(x)

A predicate which returns 0 if x is a list, and 1 if x is a non-list (atomic) value.

```
KeRF> atom `A"
1
KeRF> atom "A String"
0
KeRF> atom 37
1
KeRF> atom -0.2
1
KeRF> atom 2015.03.31
1
KeRF> atom null
1
KeRF> atom [2, 5, 16]
0
KeRF> atom a: 45, b: 76
1
```

9.11 avg - Average

avg(x)

Calculate the arithmetic mean of the elements of a list x. Equivalent to (sum x)/count_nonnull x.

```
KeRF> avg 3 7 12.5 9
7.875
```

9.12 `between` - Between?

`between(x, y)`

Predicate which returns 1 if `x` is between the first two elements of the list `y`. Equivalent to `(x >= y[0]) & (x <= y[1])`.

```
KeRF> between(2 5 17, 3 10)
[0, 1, 0]
```

Be careful- `between` will always fail if `y` is not a list or does not have the correct length:

```
KeRF> between(2 5 17, 3)
[0, 0, 0]

KeRF> 3[1]
NAN
```

9.13 `btree` - BTree

`btree(x)`

Equivalent to `indexed(x)`.

9.14 `bucketed` - Bucket Values

`bucketed(x, y)`

Equivalent to `floor (<<y) * x / count y`.

9.15 `build_table_from_csv` - Build Table From CSV File

`build_table_from_csv(tableFile, csvFile, fields, n)`

Load a Comma-Separated Value file and create a table which is serialized on disk. If the table file already exists, this operation will append. `tableFile` and `csvFile` are filenames on disk. `fields` is a string which indicates the expected datatype of each column in the CSV file- see `read_table_from_delimited_file` for the supported column types and their symbols. `n` indicates how many rows of the file are treated as column headers- generally 0 or 1.

This routine will create a striped table if the table does not already exist on disk, and otherwise will operate on a normal table as created by `write_to_path`. See **Striped Files**.

9.16 `build_table_from_fixed_file` - Build Table From Fixed-Width File

`build_table_from_fixed_file(tableFile, fixedFile, attributes)`

Load a file with fixed-width columns and create a table which is serialized on disk. If the table file already exists, this operation will append. `tableFile` and `fixedFile` are filenames on disk. `attributes` is a map specifying the format of the fixed-width file, as in `read_table_from_fixed_file`.

This routine will create a striped table if the table does not already exist on disk, and otherwise will operate on a normal table as created by `write_to_path`. See **Striped Files**.

9.17 `build_table_from_psv` - Build Table From PSV File

```
build_table_from_psv(tableFile, psvFile, fields, n)
```

Equivalent to `build_table_from_csv` which processes pipe-delimited fields (`|`) instead of comma-delimited fields.

9.18 `build_table_from_tsv` - Build Table From TSV File

```
build_table_from_tsv(tableFile, tsvFile, fields, n)
```

Equivalent to `build_table_from_csv` which processes tab-delimited fields instead of comma-delimited fields.

9.19 `car` - Contents of Address Register

```
car(x)
```

Select the first element of the list `x`. Atomic types are unaffected by this operation. Equivalent to `first(x)`. `car` is a reference to the Lisp primitive of the same name, which selected the first element of a pair.

```
KeRF> car 32 83 90
32
KeRF> car 409
409
KeRF> nil = car []
1
```

9.20 `cdr` - Contents of Data Register

```
cdr(x)
```

Select all the elements of the list `x` except for the first. Atomic types are unaffected by this operation. Equivalent to `drop(1, x)`. `cdr` is a reference to the Lisp primitive of the same name, which selected the second element of a pair.

```
KeRF> cdr 32 83 90
[83, 90]
KeRF> cdr 409
409
```

9.21 `ceil` - Ceiling

```
ceil(x)
```

Compute the smallest integer following a number `x`. Atomic.

```
KeRF> ceil -3.2 0.4 0.9 1.1
[-3, 1, 1, 2]
```

Taking the ceiling of a string or char converts it to uppercase:

```
KeRF> ceil "Hello, World!"
"HELLO, WORLD!"
```

9.22 char - Cast to Char

`char(x)`

Cast a number or list `x` to a char or string, respectively.

```
KeRF> char 65
~ "A"
KeRF> char 66.7
~ "B"
KeRF> char 72 101 108 108 111 44 32 75 101 82 70 33
"Hello, KeRF!"
```

9.23 close_socket - Close Socket

`close_socket(handle)`

Given a socket handle as obtained with `open_socket`, close the connection. See **Network I/O**.

9.24 combinations - Combinations

`combinations(x, n)`

`combinations(x, n, repeats)`

Produce a list of all the distinct subsets of `x` which contain `n` elements. Normally this will operate on the unique elements of `x`, but if `repeats` is truthy all elements will be preserved:

```
KeRF> combinations(6 7 8, 2)
[[6, 7],
 [6, 8],
 [7, 8]]
KeRF> combinations(3 5 5 7, 3)
[[3, 5, 7]]
KeRF> combinations(3 5 5 7, 3, 1)
[[3, 5, 5],
 [3, 5, 7],
 [3, 5, 7],
 [5, 5, 7]]
```

If `x` is a map, operate on its keys:

```
KeRF> combinations({foo: 1, bar: 2, quux: 3}, 2)
[["foo", "bar"],
 ["foo", "quux"],
 ["bar", "quux"]]
```

9.25 cos - Cosine

`cos(x)`

Calculate the cosine of `x`, expressed in radians. Atomic. The results of `cos` will always be floating point values.

```
KeRF> cos 3.14159 1 -20
      [-1, 0.540302, 0.408082]
KeRF> acos cos 3.14159 1 -20
      [3.14159, 1, 1.15044]
```

9.26 cosh - Hyperbolic Cosine

`cosh(x)`

Calculate the hyperbolic cosine of `x`, expressed in radians. Atomic. The results of `cosh` will always be floating point values.

```
KeRF> cosh 3.14159 1 -20
      [11.5919, 1.54308, 2.42583e+08]
```

9.27 count - Count

`count(x)`

Equivalent to `len(x)`.

9.28 count_nonnull - Count Non-Nulls

`count_nonnull(x)`

Determine the number of elements in `x` which are not null. Equivalent to `sum not isnull x`.

```
KeRF> count_nonnull 1 2 3
      3
KeRF> count_nonnull [nan, null, 45]
      1
```

9.29 count_null - Count Nulls

`count_null(x)`

Determine the number of elements in `x` which are null. Equivalent to `sum isnull x`.

```
KeRF> count_null 1 2 3
      0
KeRF> count_null [nan, null, 45]
      2
```

9.30 cross - Cartesian Product

`cross(x, y)`

Pair up each element of x with each element of y. If x or y is a map, use the values:

```
KeRF> cross(5, 3 4 6)
[[5, 3],
 [5, 4],
 [5, 6]]
KeRF> cross(1 2 3, 4 5)
[[1, 4],
 [1, 5],
 [2, 4],
 [2, 5],
 [3, 4],
 [3, 5]]
KeRF> cross(1 2 3, {a:4, b:5})
[[1, 4],
 [1, 5],
 [2, 4],
 [2, 5],
 [3, 4],
 [3, 5]]
```

9.31 deal - Deal

`deal(count, x)`

Select count random elements from the list x without repeats.

```
KeRF> deal(5, 21 31 41 51 61 71 81)
[81, 21, 61, 51, 41]
KeRF> deal(4, "ABCD")
"CDBA"
```

If x is a map or table, select keys.

```
KeRF> deal(3, {a: 99, b: 22, c: 41, d: 55})
["d", "a", "b"]
KeRF> deal(2, {{a: 99 98, b: 21 22, c: 33 34}})
["c", "b"]
```

If x is a number, select from range(x):

```
KeRF> deal(10, 10)
[7, 4, 0, 2, 9, 8, 1, 6, 3, 5]
KeRF> deal(5, 5.0)
[2, 3, 1, 0, 4]
```

9.32 descend - Descending Indices

`descend(x)`

For a list `x`, generate a list of indices into `x` in descending order of the values of `x`.

```
KeRF> t:5 2 3 1
      [5, 2, 3, 1]
KeRF> descend t
      [0, 2, 1, 3]
KeRF> t[descend t]
      [5, 3, 2, 1]
```

Strings are sorted in lexicographic order:

```
KeRF> descend ["Orange", "Apple", "Pear", "Aardvark", "A"]
      [2, 0, 1, 3, 4]
```

When applied to a map, `descend` will sort the keys by their values and produce a list:

```
KeRF> ascend {"A":2, "B":9, "C":0}
      ["B", "A", "C"]
```

The symbol `>` is equivalent to `descend` when used as a unary operator:

```
KeRF> >5 2 3 1
      [0, 2, 1, 3]
```

9.33 dir_ls - Directory Listing

`dir_ls(path)`

`dir_ls(path, full)`

List the files and directories at a filesystem path. If `full` is provided and truthy, list complete paths to the elements of the directory. Otherwise, list only the base names. See **File I/O**.

9.34 display - Display

`display(x)`

Print a display representation of data to standard output. See **General I/O**.

9.35 `distinct` - Distinct Values

`distinct(x)`

Select the first instance of each item in a list `x`. Atomic types are unaffected by this operation.

```
KeRF> distinct "BANANA"
"BAN"
KeRF> distinct 2 3 3 5 3 4 5
[2, 3, 5, 4]
```

The symbol `%` is equivalent to `distinct` when used as a unary operator:

```
KeRF> %"BANANA"
"BAN"
```

9.36 `divide` - Divide

`divide(x, y)`

Divide `x` by `y`. Fully atomic. The results of `divide` will always be floating point values.

```
KeRF> divide(3, 5)
0.6
KeRF> divide(-1, 0)
-inf
KeRF> divide(3, 2 4 5 0)
[1.5, 0.75, 0.6, inf]
KeRF> divide(1 3 4 0, 9)
[0.111111, 0.333333, 0.444444, 0.0]
KeRF> divide(10 5 3, 7 9 3)
[1.42857, 0.555556, 1.0]
```

The symbol `/` is equivalent to `divide` when used as a binary operator:

```
KeRF> 10 5 3 / 7 9 3
[1.42857, 0.555556, 1.0]
```


9.37 dload - Dynamic Library Load

`dload(filename, function, argcount)`

Load a dynamic library function and return a Kerf function which can be invoked to call into it. `filename` is the name of the library, `function` is the name of the function and `argcount` is the number of arguments the function takes. The current implementation of `dload` permits a maximum `argcount` of 8.

Let's look at a very simple C function which can be called from Kerf:

```
#include <stdio.h>
#include "kerf_api.h"

KERF foreign_function_example(KERF argument) {
    int64_t value = argument->i;
    printf("Hello from C! You gave me a %d.\n", (int)value);
    return 0;
}
```

```
KerF> f: dload("example.dylib", "foreign_function_example", 1)
      {OBJECT:foreign_function_example}
KerF> f(42)
Hello from C! You gave me a 42.
```

For more information about writing dynamic libraries for use in Kerf, see **FFI**.

9.38 dotp - Dot Product

`dotp(x, y)`

Calculate the dot product (or *scalar product*) of the vectors `x` and `y`. Equivalent to `sum x*y`.

```
KerF> dotp(1 2 3, 1 2 5)
20
```

9.39 drop - Drop Elements

`drop(count, x)`

Remove `count` elements from the beginning of the list `x`. Atomic types are unaffected by this operation.

```
KerF> drop(3, "My Hero")
"Hero"
KerF> drop(5, 9 2 0)
[]
KerF> drop(3, 5)
5
```

The symbol `_` is equivalent to `drop` when used as a binary operator:

```
KerF> 3 _ "My Hero"
"Hero"
```

9.40 enlist - Enlist Element

`enlist(x)`

Wrap any element `x` in a list.

```
KeRF> enlist "A"  
["A"]  
KeRF> enlist 22 33  
[[22, 33]]
```

9.41 enum - Enumeration

`enum(x)`

Equivalent to `hashed(x)`.

9.42 enumerate - Enumerate Items

`enumerate(x)`

If `x` is a number, generate a range of integers from 0 up to but not including `x`. Equivalent to `til(x)`.

```
KeRF> enumerate 0  
INT[]  
KeRF> enumerate 3  
[0, 1, 2]  
KeRF> enumerate 5.3  
[0, 1, 2, 3, 4]
```

If `x` is a map, extract its keys.

```
KeRF> enumerate b:43, a:999  
["b", "a"]
```

If `x` is a list, generate the Cartesian Product over the ranges of each element of `x`. This operation is sometimes called *odometer*, for the way the resulting lists resemble the rolling digits of a car's odometer.

```
KeRF> enumerate 2 3  
[[0, 0],  
 [0, 1],  
 [0, 2],  
 [1, 0],  
 [1, 1],  
 [1, 2]]  
KeRF> enumerate 2 2 2  
[[0, 0, 0],  
 [0, 0, 1],  
 [0, 1, 0],  
 [0, 1, 1],  
 [1, 0, 0],  
 [1, 0, 1],  
 [1, 1, 0],  
 [1, 1, 1]]
```

The symbol `^` is equivalent to `enumerate` when used as a unary operator:

```
KeRF> ^9
[0, 1, 2, 3, 4, 5, 6, 7, 8]
```

9.43 `equal` - Equal?

`equal(x, y)`

A predicate which returns 1 if `x` is equal to `y`. Equivalent to `equals(x)`. Fully atomic.

```
KeRF> equal(5, 13)
0
KeRF> equal(5, 5 13)
[1, 0]
KeRF> equal(5 13, 5 13)
[1, 1]
KeRF> equal(.1, .1000000000000001)
0
KeRF> equal(nan, nan)
1
```

The symbols `=` and `==` are equivalent to `equal` when used as binary operators:

```
KeRF> 3 == 1 3 5
[0, 1, 0]
```

9.44 `equals` - Equals?

`equals(x, y)`

Equivalent to `equal(x)`.

9.45 `erf` - Error Function

`erf(x)`

Compute the Gauss error function of `x`. Atomic.

```
KeRF> erf -.5 -.2 0 .2 .3 1 2
[-0.5205, -0.222703, 0, 0.222703, 0.328627, 0.842701, 0.995322]
```

9.46 `erfc` - Complementary Error Function

`erfc(x)`

Compute the complementary Gauss error function of `x`. Equivalent to `1 - erf(x)`. Atomic.

```
KeRF> erfc -.5 -.2 0 .2 .3 1 2
[1.5205, 1.2227, 1, 0.777297, 0.671373, 0.157299, 0.00467773]
```

9.47 eval - Evaluate

`eval(x)`

Evaluate a string `x` as a Kerf expression. Atomic down to strings.

```
Kerf> a
Kerf> eval(["2+3", "a: 24", "a: 999"])
[5, 24, 999]
Kerf> a
999
```

9.48 except - Except

`except(x, y)`

Remove all the elements of `y` from `x`. Equivalent to `x[which not x in y]`.

```
Kerf> except("ABCDDBEFB", "ADF")
"BCBEB"
```

If `x` is atomic, the result will be enclosed in a list:

```
Kerf> except(2, 3 4)
[2]
```

9.49 exit - Exit

`exit()`
`exit(code)`

Exit the Kerf interpreter. If a number is provided, use it as an exit code. Otherwise, exit with code 0 (successful).

```
Kerf> exit(1)
```

9.50 exp - Natural Exponential Function

`exp(x)`
`exp(x, y)`

Calculate e^x , the natural exponential function. If `y` is provided, calculate x^y . Fully atomic.

```
Kerf> exp 1 2 5
[2.71828, 7.38906, 148.413]
Kerf> exp(2, 0 1 2)
[1, 2, 4.0]
```

The symbol `**` is equivalent to `exp`:

```
Kerf> **1 2 5
[2.71828, 7.38906, 148.413]
Kerf> 2**0 1 2
[1, 2, 4.0]
```

9.51 explode - Explode

`explode(key, x)`

Violently and suddenly split the list `x` at instances of `key`. To reverse this process, use `implode`.

```
KeRF> explode(`"e", "A dream deferred")
["A dr", "am d", "f", "rr", "d"]
KeRF> explode(0, 1 1 2 0 2 0 5)
[[1, 1, 2], [2], [5]]
```

`explode` does not search for subsequences.

Splitting on a 1-length string is not the same as splitting on a character:

```
KeRF> explode("rat", "drat, that rat went splat.")
["drat, that rat went splat."]
KeRF> explode("e", "A dream deferred")
["A dream deferred"]
```

9.52 filter - Filter

`filter(f, x)`

Apply a predicate `f` to each element of `x`, and return a list of elements where `f` produces a truthy value:

```
KeRF> (not isnull) filter [3, null, 27, nil, 39]
[3, 27, 39]
KeRF> {[x] x > 100} filter 20 19 130 5 -2 200 119
[130, 200, 119]
KeRF> {[x] len(x) = 4} filter ["apple","pear","knife","lock","spider"]
["pear", "lock"]
```

If `x` is a map, `filter` will operate on the values of `x` and produce a map result:

```
KeRF> filter(isnull, {a:23,b:null,c:22,e:null})
{b:null, e:null}
KeRF> {[x] len(x) > 1} filter {foo:1 2 3, bar:4, quux:5 6}
{foo:[1, 2, 3], quux:[5, 6]}
```

For tables, prefer using SQL syntax and the more convenient `WHERE` clause.

9.53 first - First

`first(x)`

`first(x, y)`

When provided with a single argument, select the first element of the list `x`. Atomic types are unaffected by this operation.

```
KeRF> first(43 812 99 23)
43
KeRF> first(99)
99
```

When provided with two arguments, select the first x elements of y, repeating elements of y as necessary. Equivalent to take(x, y).

```
KeRF> first(2, 43 812 99 23)
[43, 812]
KeRF> first(8, 43 812 99 23)
[43, 812, 99, 23, 43, 812, 99, 23]
```

9.54 flatten - Flatten

flatten(x)

Concatenate the elements of the list x. To join elements with a delimiter, use implode.

```
KeRF> flatten(["foo", "bar", "quux"])
"foobarquux"
KeRF> flatten([2 3 4, 9 7 8, 14])
[2, 3, 4, 9, 7, 8, 14]
```

Note that flatten only removes one level of nesting. To completely flatten an arbitrarily nested structure, combine it with converge:

```
KeRF> n: [[1,2],[3,4],[5,[6,7]]];
KeRF> flatten n
[1, 2, 3, 4, 5, [6, 7]]
KeRF> flatten converge n
[1, 2, 3, 4, 5, 6, 7]
```

9.55 float - Cast to Float

float(x)

Cast x to a float. Atomic.

```
KeRF> float 0 7 15
[0, 7, 15.0]
```

When applied to a string, parse it into a number:

```
KeRF> float "97"
97.0
```

9.56 floor - Floor

floor(x)

Compute the largest integer preceding a number x. Atomic.

```
KeRF> floor -3.2 0.4 0.9 1.1
[-4, 0, 0, 1]
KeRF> int -3.2 0.4 0.9 1.1
[-3, 0, 0, 1]
```

Taking the floor of a string or char converts it to lowercase:

```
KeRF> floor "Hello, World!"  
"hello, world!"
```

The symbol `_` is equivalent to `floor` when used as a unary operator:

```
KeRF> _ 37.9 14.2  
[37, 14]
```

9.57 greater - Greater Than?

`greater(x, y)`

A predicate which returns 1 if `x` is greater than `y`. Fully atomic.

```
KeRF> greater(1 2 3, 2)  
[0, 0, 1]  
KeRF> greater([5], [[], [3], [2 9]])  
[0, 1, 0]  
KeRF> greater("apple", ["a", "aa", "banana"])  
[1, 1, 0]
```

The symbol `>` is equivalent to `greater` when used as a binary operator:

```
KeRF> 3 4 7 > 1 9 0  
[1, 0, 1]
```

9.58 greatereq - Greater or Equal?

`greatereq(x, y)`

A predicate which returns 1 if `x` is greater than or equal to `y`. Fully atomic.

```
KeRF> greatereq(1 2 3, 2)  
[0, 1, 1]
```

The symbol `>=` is equivalent to `greatereq` when used as a binary operator:

```
KeRF> 3 4 5 7 >= 1 9 5 0  
[1, 0, 1, 1]
```

9.59 has_column - Table Has Column?

`has_column(table, key)`

A predicate which returns 1 if `table` has a column with the key `key`.

```
KeRF> has_column({{a: 1 2 3; b: 4 2 1}}, "a")  
1  
KeRF> has_column({{a: 1 2 3; b: 4 2 1}}, "fictional")  
0
```

9.60 has_key - Has Key?

`has_key(x, key)`

A predicate which returns 1 if a map `x` contains the key `key`.

```
KeRF> m: {alphonse: 1, betty: 3, oscar: 99};
KeRF> has_key(m, "alphonse")
1
KeRF> has_key(m, "alphys")
0
```

If `x` is a list, return 1 if `key` is a valid index into `x`:

```
KeRF> l: 45 99 10 15;
KeRF> has_key(l, -1)
0
KeRF> has_key(l, 2)
1
KeRF> has_key(l, 2.2)
1
KeRF> l[2.2]
10
KeRF> l[-1]
NAN
```

If `x` is a table, equivalent to `has_column(x, key)`.

9.61 hash - Hash

`hash(x)`

Equivalent to `hashed(x)`.

9.62 hashed - Hashed

`hashed(x)`

Create a list containing the elements of `x`, with hashmap-backed local interning. Interning will minimize the storage consumed by values which occur frequently and permit much more efficient sorting.

```
KeRF> data: rand(10000, ["apple", "pear", "banana"]);
KeRF> write_to_path("a.data", data);
KeRF> write_to_path("b.data", #data);
KeRF> shell "wc -c a.data"
[" 1048576 a.data"]
KeRF> shell "wc -c b.data"
[" 262144 b.data"]
```

The symbol `#` is equivalent to `hashed` when used as a unary operator:

```
KeRF> #["a", "b", "a"]
#["a", "b", "a"]
```


9.63 **ident** - Identity

`ident(x)`

Unary identity function. Returns x unchanged.

```
KeRF> ident 42
42
```

The symbol `:` is equivalent to `ident` when used as a unary operator:

```
KeRF> :42
42
```

9.64 **ifnull** - If Null?

`ifnull(x)`

Equivalent to `isnull(x)`.

9.65 **implode** - Implode

`implode(key, x)`

Violently and suddenly join the elements of the list `x` intercalated with `key`. To reverse this process, use `explode`.

```
KeRF> implode(" and_", ["BIFF", "BOOM", "POW"])
"BIFF_and_BOOM_and_POW"
KeRF> implode(23, 10 4 3 15)
[10, 23, 4, 23, 3, 23, 15]
```

9.66 **in** - In?

`in(key, x)`

A predicate which returns 1 if each key is an element of `x`. Atomic over `key`.

```
KeRF> in(3, 8 7 3 2)
[1]
KeRF> in(3 4, 8 7 3 2)
[1, 0]
KeRF> in("a", "cassiopeia")
[1]
```

9.67 **index** - Index

`index(x)`

Equivalent to `indexed(x)`.

9.68 indexed - Indexed

`indexed(x)`

Create an indexed version of a list `x`. This constructs an associated B-Tree, permitting faster searches and range queries. Do not use `indexed` if you know the list must always be ascending. The command `sort_debug` can be used to determine whether Kerf thinks a list is already sorted.

```
KerF> indexed 3 7 0 5 2
=[3, 7, 0, 5, 2]
```

The symbol `=` is equivalent to `indexed` when used as a unary operator:

```
KerF> =3 2
=[3, 2]
```

9.69 int - Cast to Int

`int(x)`

Cast `x` to an int, truncating. Atomic.

```
KerF> int 33.6 -12.5 4 nan
[33, -12, 4, NAN]
KerF> floor 33.6 -12.5 4 nan
[33, -13, 4, NAN]
```

When applied to a string, parse it into an integer:

```
KerF> int ["1337", "47.2", ""]
[1337, 47, 0]
```

9.70 intersect - Set Intersection

`intersect(x, y)`

Find unique items contained in both `x` and `y`. Equivalent to `distinct(x)[which distinct(x) in y]`.

```
KerF> intersect(4, 4 5 6)
[4]
KerF> intersect(3 4, 1 2 3 4 5 6)
[3, 4]
KerF> intersect("ABD", "BCBD")
"BD"
```

9.71 isnull - Is Null?

`isnull(x)`

A predicate which returns 1 if `x` is null. Atomic.

```
KerF> isnull([(), nan, 2, -3.7, [], {a:5}])
[1, 1, 0, 0, [], {a:0}]
```

9.72 join - Join

`join(x, y)`

Form a list by catenating x and y.

```
KerF> join(1, 2)
[1, 2]
KerF> join(2, 3 4)
[2, 3, 4]
KerF> join(2 3, 4 5)
[2, 3, 4, 5]
KerF> join(2 3, 4)
[2, 3, 4]
KerF> join(2 3, "ABC")
[2, 3, `"A", `"B", `"C"]
KerF> join({a:23, b:24}, {b:99})
[{a:23, b:24}, {b:99}]
```

The symbol # is equivalent to join when used as a binary operator:

```
KerF> 2 3 # 9
[2, 3, 9]
KerF> "foo" #(rep 23)#!"
"foo 23!"
```

9.73 json_from_kerf - Convert Kerf to JSON

`json_from_kerf(x)`

Convert a Kerf data structure x into a JSON (IETF RFC-4627) string.

```
KerF> json_from_kerf({a: 45, b: [1, 3, 5.0]})
"{\"a\":45,\"b\":[1,3,5]}"
KerF> json_from_kerf({{a: 1 2 3, b: 4 5 6}})
"{\"a\":[1,2,3],\"b\":[4,5,6],\"is_json_table\":[1]}"
```

9.74 kerf_from_json - Convert JSON to Kerf

kerf_from_json(string)

Convert a JSON (IETF RFC-4627) string into a Kerf data structure. Note that booleans become the numbers 1 and 0 during this conversion process. Kerf-generated JSON strings generally contain the metadata necessary to round-trip without information loss, but JSON strings produced by another program may not.

```
KerF> kerf_from_json("[23, 45, 9]")
[23, 45, 9]
KerF> kerf_from_json("[true, false]")
[1, 0]
KerF> kerf_from_json("{\"a\":[1,2,3],\"b\":[4,5,6]}")
a:[1, 2, 3], b:[4, 5, 6]
KerF> kerf_from_json("{\"a\":[1,2,3],\"b\":[4,5,6],\"is_json_table\":[1]}")
```

a	b
1	4
2	5
3	6

9.75 kerf_type - Type Code

kerf_type(x)

Obtain a numeric typecode from a Kerf value.

```
KerF> kerf_type 45.0
3
```

Type	Example	kerf_type_name	kerf_type
Timestamp Vector	[2000.01.01]	stamp vector	-4
Float Vector	[0.1]	float vector	-3
Integer Vector	[1]	integer vector	-2
Character Vector	"A"	character vector	-1
Function	{[x] 1+x}	function	0
Character	`A	character	1
Integer	1	integer	2
Float	0.1	float	3
Timestamp	2000.01.01	stamp	4
Null	()	null	5
List	[]	list	6
Map	{a:1}	map	7
Enumeration	enum ["a"]	enum	8
Index	index [1,2]	sort	9
Table	{{a:1}}	table	10
Atlas	atlas {a:1}	atlas	11

Kerf types

9.76 `kerf_type_name` - Type Name

`kerf_type_name(x)`

Obtain a human-readable type name string from a Kerf value. See `kerf_type`.

```
KerF> kerf_type_name "Text"  
"character vector"
```

9.77 `last` - Last

`last(x)`

`last(count, x)`

When provided with a single argument, select the last element of the list `x`. Atomic types are unaffected by this operation.

```
KerF> last(43 812 99 23)  
23  
KerF> last(99)  
99
```

When provided with two arguments, select the last count elements of `x`, repeating elements of `x` as necessary. Equivalent to `take(-count, x)`.

```
KerF> last(2, 43 812 99 23)  
[99, 23]  
KerF> last(7, 43 812 99 23)  
[812, 99, 23, 43, 812, 99, 23]
```

9.78 `left_join` - Left Join

`left_join(x, y, z)`

Perform a left join of the tables `x` and `y` on the column `z`. See **Joins**.

9.79 `len` - Length

`len(x)`

Determine the number of elements in `x`. Equivalent to `count(x)`. Atomic elements have a count of 1.

```
KerF> len 4 7 9  
3  
KerF> len [4 7 9, 23 32]  
2  
KerF> len 5  
1  
KerF> len {a:23, b:45}  
1
```

9.80 less - Less Than?

`less(x, y)`

A predicate which returns 1 if `x` is less than `y`. Fully atomic.

```
KeRF> less(1 2 3, 2)
[1, 0, 0]
KeRF> less([5], [[], [3], [2 9]])
[1, 0, 1]
KeRF> less("apple", ["a", "aa", "banana"])
[0, 0, 1]
```

The symbol `<` is equivalent to `less` when used as a binary operator:

```
KeRF> 3 4 7 < 1 9 0
[0, 1, 0]
```

9.81 lesseq - Less or Equal?

`lesseq(x, y)`

A predicate which returns 1 if `x` is less than or equal to `x`. Fully atomic.

```
KeRF> lesseq(1 2 3, 2)
[1, 1, 0]
KeRF> lesseq([5], [[], [3], [5], [2 9]])
[1, 0, 0, 1]
KeRF> lesseq("apple", ["a", "aa", "apple", "banana"])
[0, 0, 1, 1]
```

The symbol `<=` is equivalent to `lesseq` when used as a binary operator:

```
KeRF> 3 4 1 7 <= 1 9 1 0
[0, 1, 1, 0]
```

9.82 lg - Base 2 Logarithm

`lg(x)`

Calculate $\log_2(x)$. Equivalent to $\log(2, x)$. Atomic.

```
KeRF> lg 128 512 37
[7, 9, 5.20945]
```

9.83 lines - Lines From File

`lines(filename)`
`lines(filename, n)`

Load lines from `filename` into a list of strings. If `n` is present, limit loading to `n` lines. See **File I/O**.

9.84 ln - Natural Logarithm

ln(x)

Calculate $\log_e(x)$. Atomic.

```
KeRF> ln 2 3 10 37
[0.693147, 1.09861, 2.30259, 3.61092]
```

9.85 load - Load Source

load(filename)

Load and run Kerf source from a file. Given an example file:

```
// comment
a: 7+range 10
b: range 10
a * b
```

Loading the file from the Repl:

```
KeRF> load("manual/example.kerf")
KeRF> a
[7, 8, 9, 10, 11, 12, 13, 14, 15, 16]
KeRF> b
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

Note that the value of the raw expression `a * b` is not printed when you load a file. If this is desired, use `display` in the script.

9.86 log - Logarithm

log(x)

log(x, y)

Calculate the base x logarithm of y. If only one argument is provided, $\log(10, x)$ is assumed. Fully atomic.

```
KeRF> log 3 8 10 16 100
[0.477121, 0.90309, 1, 1.20412, 2.0]
KeRF> log(2, 3 8 10 16 100)
[1.58496, 3, 3.32193, 4, 6.64386]
KeRF> log(2 3 4, 8)
[3, 1.89279, 1.5]
```

9.87 lsq - Least Squares Solution

lsq(A, B)

Solve $Ax = B$ for x , where A is a matrix and B is a matrix or vector.

```
KeRF> lsq([1 2;4 4], [3 4])
[[1, 0.5]]
KeRF> lsq([0 .5;2 0], [0 1;1 0])
[[2, 0.0],
 [0, 0.5]]
```

The symbol \ is equivalent to lsq when used as a binary operator:

```
KeRF> [0 .5;2 0]\[0 1;1 0]
[[2, 0.0],
 [0, 0.5]]
```

9.88 map - Make Map

map(keys, values)

Make a map from a list of keys and a list of values. These lists must be the same length.

```
KeRF> map("ABC", 44 18 790)
{"A":44, "B":18, "C":790}
```

The symbol ! is equivalent to map when used as a binary operator:

```
KeRF> "ABC" ! 44 18 790
{"A":44, "B":18, "C":790}
```

9.89 match - Match?

match(x, y)

A predicate which returns 1 if x is identical to y. While equals compares atoms, match is not atomic and compares entire values.

```
KeRF> match(5, 3 5 7)
0
KeRF> match(3 5 7, 3 5 7)
1
```

The symbol ~ is equivalent to match when used as a binary operator:

```
KeRF> 3 5 7 ~ 3 5 7
1
```


9.90 mavg - Moving Average

mavg(x, y)

For a series of sliding windows of size x ending at each element of the list y, find the arithmetic mean of valid (not nan and in range of the list) elements. Equivalent to msum(x, y)/mcount(x, y).

```
KeRF> mavg(3, 1 2 4 2 1 nan 0 4 6 7)
[1, 1.5, 2.33333, 2.66667, 2.33333, 1.5, 0.5, 2, 3.33333, 5.66667]
```

9.91 max - Maximum

max(x)

Find the maximum element of x. Roughly equivalent to last sort x, but much more efficient.

```
KeRF> max 7 0 15 -1 8
15
KeRF> max 0 -inf -5 nan inf
inf
```

9.92 maxes - Maximums

maxes(x, y)

Find the maximum of x and y. Fully atomic.

```
KeRF> maxes(1 3 7 2 0 1, -4 3 20 1 9 4)
[1, 3, 20, 2, 9, 4]
KeRF> maxes(8, -4 3 20 1 9 4)
[8, 8, 20, 8, 9, 8]
```

The symbol | is equivalent to maxes when used as a binary operator:

```
KeRF> -4 3 20 1 9 4 | 15
[15, 15, 20, 15, 15, 15]
```

9.93 mcount - Moving Count

mcount(x, y)

For a series of sliding windows of size x ending at each element of the list y, count the number of valid (not nan and in range of the list) elements. Equivalent to msum(x, not isnull y).

```
KeRF> mcount(3, 1 2 3 nan 4 5 nan nan nan)
[1, 2, 3, 2, 2, 2, 2, 1, 0]
```

9.94 median - Median

`median(x)`

Select the median of a list of numbers x:

```
KeRF> median 1 3 19 7 4 2
3.5
```

9.95 meta_table - Meta Table

`meta_table(table)`

Produce a table containing debugging metadata about some some table:

```
KeRF> meta_table {{a: 1 2 3, b: 3.0 17 4}}
```

column	type	type_name	is_ascending	is_disk
a	-2	integer vector	1	0
b	-3	float vector	0	0

9.96 min - Minimum

`min(x)`

Find the minimum element of x. Roughly equivalent to `first sort`, but much more efficient.

```
KeRF> min 7 0 15 -1 8
-1
KeRF> min 0 -inf -5 nan inf
-inf
```

9.97 mins - Minimums

`mins(x, y)`

Find the minimum of x and y. Fully atomic.

```
KeRF> mins(1 3 7 2 0 1, -4 3 20 1 9 4)
[-4, 3, 7, 1, 0, 1]
KeRF> mins(8, -4 3 20 1 9 4)
[-4, 3, 8, 1, 8, 4]
```

The symbol `&` is equivalent to `mins` when used as a binary operator:

```
KeRF> -4 3 20 1 9 4 & 15
[-4, 3, 15, 1, 9, 4]
```

9.98 minus - Minus

`minus(x, y)`

Calculate the difference of x and y. Fully atomic.

```
KeRF> minus(3, 5)
-2
KeRF> minus(3, 9 15 -7)
[-6, -12, 10]
KeRF> minus(9 15 -7, 3)
[6, 12, -10]
KeRF> minus(9 15 -7, 1 3 5)
[8, 12, -12]
```

The symbol `-` is equivalent to `minus` when used as a binary operator:

```
KeRF> 2 4 3 - 9
[-7, -5, -6]
```

9.99 minv - Matrix Inverse

`minv(x)`

Calculate the inverse of a matrix x.

```
KeRF> minv([1 2;3 4])
[[-2,    1 ],
 [ 1.5, -0.5]]
```

9.100 mmax - Moving Maximum

`mmax(x, y)`

For a series of sliding windows of size x ending at each element of the list y, find the maximum element. Equivalent to `(x-1)` or `mapback converge y`.

```
KeRF> mmax(3, 0 1 0 2 0 1 0)
[0, 1, 1, 2, 2, 2, 1]
```

9.101 mmin - Moving Minimum

`mmin(x, y)`

For a series of sliding windows of size x ending at each element of the list y, find the minimum element. Equivalent to `(x-1)` and `mapback converge y`.

```
KeRF> mmin(3, 4 0 3 0 2 0 4 5 6)
[4, 0, 0, 0, 0, 0, 0, 0, 4]
```

9.102 mmul - Matrix Multiply

`mmul(x, y)`

Multiply the matrix or vector `x` by the matrix or vector `y`. Equivalent to `x dotp mapleft y`.

```
KeRF> mmul([1 2;3 4], [5 6])
[[15, 18],
 [35, 42]]
KeRF> mmul([1 2;3 4], [0 1;1 0])
[[2, 1],
 [4, 3]]
```

9.103 mod - Modulus

`mod(x, y)`

Calculate `x` modulo `y`. Equivalent to `x - y * floor(x/y)`. Left-atomic.

```
KeRF> mod(0 1 2 3 4 5 6 7, 3)
[0, 1, 2, 0, 1, 2, 0, 1]
KeRF> mod(-4 -3 -2 -1 0 1 2, 2)
[0, 1, 0, 1, 0, 1, 0]
```

The symbol `%` is equivalent to `mod` when used as a binary operator:

```
KeRF> 0 1 2 3 4 5 6 7 % 3
[0, 1, 2, 0, 1, 2, 0, 1]
```

9.104 msum - Moving Sum

`msum(x, y)`

Calculate a series of sums of each element in a list `y` and up to the `x` previous values, ignoring nans and nonexistent values.

```
KeRF> msum(2, 10 20 30 40)
[10, 30, 50, 70]
KeRF> msum(2, 1 2 2 nan 1 2)
[1, 3, 4, 2, 1, 3.0]
KeRF> msum(3, 10 10 14 10 25 10 Nan 10)
[10, 20, 34, 34, 49, 45, 35, 20.0]
```

`msum(1, y)` can be used to remove nan from data:

```
KeRF> msum(1, 4 2 1 nan 2)
[4, 2, 1, 0, 2.0]
```

9.105 negate - Negate

`negate(x)`

Reverse the sign of a number `x`. Equivalent to `-1 * x`. Atomic.

```
KeRF> negate 2 4 -77  
[-2, -4, 77]
```

The symbol `-` is equivalent to `negate` when used as a binary operator:

```
KeRF> -(2 4 -77)  
[-2, -4, 77]
```

9.106 negative - Negative

`negative(x)`

Equivalent to `negate(x)`.

9.107 ngram - N-Gram

`ngram(n, x)`

Break a list `x` into non-overlapping subsequences of length `n`.
Equivalent to `split(range(0,count(x),n), x)`.

```
KeRF> ngram(3, "Prisencolinensinainciusol")  
["Pri", "sen", "col", "ine", "nsi", "nai", "nci", "uso", "l"]  
KeRF> ngram(2, 34 12 44 19 29 90)  
[[34, 12], [44, 19], [29, 90]]
```

9.108 not - Logical Not

`not(x)`

Calculate the logical *NOT* of `x`. Atomic.

```
KeRF> not(1 0)  
[0, 1]  
KeRF> not([0, -4, 9, nan, []])  
[1, 0, 0, 0, []]
```

The symbol `!` is equivalent to `not` when used as a unary operator:

```
KeRF> !1 0 8  
[0, 1, 0]
```

9.109 `noteq` - Not Equal?

`noteq(x, y)`

A predicate which returns 1 if x is not equal to y. Equivalent to `not equals(x)`. Fully atomic.

```
KeRF> noteq(5, 13)
1
KeRF> noteq(5, 5 13)
[0, 1]
KeRF> noteq(5 13, 5 13)
[0, 0]
KeRF> noteq(.1, .1000000000000001)
1
KeRF> noteq(nan, nan)
0
```

The symbols `!=` and `<>` are equivalent to `noteq` when used as binary operators:

```
KeRF> 3 != 1 3 5
[1, 0, 1]
```

9.110 `now` - Current DateTime

`now()`

Return a stamp containing the current date and time in UTC.

```
KeRF> now()
2015.10.31T21:14:09.018
```

9.111 `now_date` - Current Date

`now_date()`

Return a stamp containing the current date only in UTC.

```
KeRF> now_date()
2015.10.31
```

9.112 `now_time` - Current Time

`now_time()`

Return a stamp containing the current time only in UTC.

```
KeRF> now_time()
21:14:09.018
```

9.113 `open_socket` - Open Socket

`open_socket(host, port)`

Establish a connection to a remote Kerf instance at hostname `host` and listening on `port` and return a connection handle. Both `host` and `port` must be strings. See **Network I/O**.

9.114 `open_table` - Open Table

`open_table(filename)`

Load a serialized table from the binary file `filename`. See **File I/O**.

9.115 `or` - Logical OR

`or(x, y)`

Calculate the logical *OR* of `x` and `y`. This operation is equivalent to the primitive function `max`. Fully atomic.

```
Kerf> or(1 1 0 0, 1 0 1 0)
      [1, 1, 1, 0]
Kerf> or(1 2 3 4, 0 -4 9 0)
      [1, 2, 9, 4]
```

The symbol `|` is equivalent to `or` when used as a binary operator:

```
Kerf> 1 1 0 0 | 1 0 1 0
      [1, 1, 1, 0]
```

9.116 `order` - Order

`order(x)`

Generate a list of indices showing the relative ascending order of items in the list `x`. Equivalent to `<<x`.

```
Kerf> order "ABCEDF"
      [0, 1, 2, 4, 3, 5]
Kerf> order 2 4 1 9
      [1, 2, 0, 3]
Kerf> <2 4 1 9
      [2, 0, 1, 3]
Kerf> <<2 4 1 9
      [1, 2, 0, 3]
```

9.117 `out` - Output

`out(x)`

Print a string `x` to standard output. See **General I/O**.

9.118 part - Partition

part(x)

Produce a map from unique elements of a list x to lists of the indices at which these elements could originally be found.

```
KeRF> part 3 5 7 7 5
      {3:[0], 5:[1, 4], 7:[2, 3]}
KeRF> part ["apple", "frog", "frog", "kumquat"]
      {apple:[0], frog:[1, 2], kumquat:[3]}
```

part does not affect atomic types:

```
KeRF> part {a: 23 45, b: 9}
      {a:[23, 45], b:9}
KeRF> part 23
      23
```

The symbol & is equivalent to part when used as a unary operator:

```
KeRF> &2 2 1 2
      {2:[0, 1, 3], 1:[2]}
```

9.119 permutations - Permutations

permutations(x)

permutations(x, repeats)

Generate a list of all possible orderings of the elements of x. Normally this will operate on the unique elements of x, but if repeats is truthy all elements will be preserved:

```
KeRF> permutations(1 2 2)
      [[1, 2, 2],
       [2, 1, 2],
       [2, 2, 1]]
KeRF> permutations(1 2 2, 1)
      [[1, 2, 2],
       [1, 2, 2],
       [2, 1, 2],
       [2, 2, 1],
       [2, 1, 2],
       [2, 2, 1]]
```

If x is a map, operate on its keys:

```
KeRF> permutations({foo: 27, bar: 38})
      [{"foo", "bar"},
       {"bar", "foo"}]
```

9.120 plus - Plus

plus(x, y)

Equivalent to add.

9.121 pow - Exponentiation

```
pow(x)
pow(x, y)
```

Equivalent to `exp`.

9.122 powerset - Power Set

```
powerset(x)
```

Produce a list of all possible sublists of `x`. If `x` is a map, generate a powerset of its keys:

```
KeRF> powerset 45 67 33
[INT[],
 [33],
 [67],
 [67, 33],
 [45],
 [45, 33],
 [45, 67],
 [45, 67, 33]]
KeRF> powerset({foo: 23, bar: 94})
[[],
 ["bar"],
 ["foo"],
 ["foo", "bar"]]
```

9.123 rand - Random Numbers

```
rand()
rand(x)
rand(x, y)
```

Generate a random vector of `x` numbers from 0 up to but not including `y`.

```
KeRF> rand(10, 3)
[0, 2, 1, 2, 1, 2, 1, 2, 0, 2]
KeRF> rand(5, 3.0)
[0.74465, 1.72491, 0.79121, 2.53097, 0.573115]
```

If `y` is a list, select random elements from `y`.

```
KeRF> rand(6, "ABC")
"CBCBBA"
```

If y is not provided, generate a single random number from 0 up to but not including x. As above, if x is a list, choose a single random element.

```
KeRF> rand(10)
1
KeRF> rand(10)
8
KeRF> rand(10.0)
6.77151
KeRF> rand(10.0)
0.401473
KeRF> rand("ABCDE")
`"B"
```

If rand is given no arguments, generate a single random float from 0 up to but not including 1.

```
KeRF> rand()
0.389022
```

The symbol ? is equivalent to rand when used as a binary operator:

```
KeRF> 5?2
[0, 0, 1, 0, 0]
```

9.124 range - Range

```
range(x)
range(x, y)
range(x, y, z)
```

If range is provided with one argument, generate a vector of integers from 0 up to but not including x:

```
KeRF> range 5
[0, 1, 2, 3, 4]
```

If range is provided with two arguments, generate a vector of numbers from x up to but not including y, spaced 1 apart:

```
KeRF> range(10, 15)
[10, 11, 12, 13, 14]
KeRF> range(10.5, 16.5)
[10.5, 11.5, 12.5, 13.5, 14.5, 15.5]
```

If range is provided with three arguments, generate a vector of numbers from x up to but not including y, spaced z apart:

```
KeRF> range(1, 3, .3)
[1, 1.3, 1.6, 1.9, 2.2, 2.5, 2.8]
```

9.125 read_from_path - Read From Path

```
read_from_path(filename)
```

Load a serialized Kerf data structure from the binary file filename. See **File I/O**.

9.126 `read_stripped_from_path` - Read Striped File From Path

```
read_stripped_from_path(path)
```

Load a striped Kerf data structure from the directory path. See **Striped Files**.

9.127 `read_table_from_csv` - Read Table From CSV File

```
read_table_from_csv(filename, fields, n)
```

Load a Comma-Separated Value file into a table. `fields` is a string which indicates the expected datatype of each column in the CSV file- see `read_table_from_delimited_file` for the supported column types and their symbols. `n` indicates how many rows of the file are treated as column headers- generally 0 or 1.

Equivalent to `read_table_from_delimited_file(",", filename, fields, n)`.

9.128 `read_table_from_delimited_file` - Read Table From Delimited File

```
read_table_from_delimited_file(delimiter, filename, fields, n)
```

Load the contents of a text file with rows separated by newlines and fields separated by some character delimiter into a table. `fields` is a string which indicates the expected datatype of each column. `n` indicates how many rows of the file are treated as column headers- generally 0 or 1.

Symbol	Datatype
I	Integer
F	Float
S	String
E	Enumerated String
G	IETF RFC-4122 UUID
N	IP address as parsed by <code>inet_pton()</code>
Z	Custom Datetime. See <code>.Parse.strptime_format</code>
Y	Custom Times. See <code>.Parse.strptime_format2</code>
*	Skipped field

Symbols accepted as part of `fields`

Given a file like the following:

```
Language&Lines&Runtime
C&271&0.101
Java&89&0.34
Python&62&3.79
```

```
Kerf> read_table_from_delimited_file("&", "code.txt", "SIF", 1)
```

Language	Lines	Runtime
C	271	0.101
Java	89	0.34
Python	62	3.79

9.129 read_table_from_fixed_file - Read Table From Fixed-Width File

`read_table_from_fixed_file(filename, attributes)`

Load the contents of a text file with fixed-width columns. The map `attributes` specifies the details of the format:

Key	Type	Optional	Description
<code>fields</code>	String	No	As in <code>read_table_from_delimited_file</code> .
<code>widths</code>	Integer Vector	No	The width of each column in characters.
<code>line_limit</code>	Integer	Yes	The maximum number of rows to load.
<code>titles</code>	List of String	Yes	Key for each column in the resulting table.
<code>header_rows</code>	Integer	Yes	How many rows are treated as column headers.
<code>newline_separated</code>	Boolean	Yes	if false, do not expect newlines separating rows.

Settings described in `attributes`

Symbol	Datatype
R	NYSE TAQ symbol
Q	NYSE TAQ time format (HHMMSSXXX)

Additional symbols accepted as part of fields in fixed-width files

Given a file like this list of ingredients for my *famous* pizza dough:

```
Eggs      2.0 -
Flour     5.0 cups
Honey     2.0 tbs
Water     2.0 cups
Olive Oil 2.0 tbsp
Yeast     1.0 tbsp
```

```
KerF> fmt: { fields: "SFS", widths: [10, 4, 5], titles: ["Ingredient", "Amount", "Unit"] };
KerF> read_table_from_fixed_file("dough.txt", fmt)
```

Ingredient	Amount	Unit
Eggs	2.0	-
Flour	5.0	cups
Honey	2.0	tbsp
Water	2.0	cups
Olive Oil	2.0	tbsp
Yeast	1.0	tbsp

9.130 read_table_from_tsv - Read Table From TSV File

`read_table_from_tsv(filename, fields, n)`

Load a Tab-Separated Value file into a table. `fields` is a string which indicates the expected datatype of each column in the TSV file- see `read_table_from_delimited_file` for the supported column types and their symbols. `n` indicates how many rows of the file are treated as column headers- generally 0 or 1.

Equivalent to `read_table_from_delimited_file("\t", filename, fields, n)`.

9.131 rep - Output Representation

rep(x)

Convert a value x into a printable string representation. If you only wish to convert the atoms of x into strings, use string.

```
KerF> rep 45
"45"
KerF> rep 2 5 3
"[2, 5, 3]"
KerF> rep {a:4}
"{a:4}"
KerF> rep "Some text"
 "\"Some text\""
```

9.132 repeat - Repeat

repeat(n, x)

Create a list containing n copies of x. Equivalent to n take enlist x.

```
KerF> repeat(2, 5)
[5, 5]
KerF> repeat(4, "AB")
["AB", "AB", "AB", "AB"]
KerF> repeat(0, "AB")
[]
KerF> repeat(-3, "AB")
[]
```

9.133 reserved - Reserved Names

reserved()

Print and return an unsorted list of Kerf's reserved names, including reserved literals such as true.

9.134 reverse - Reverse

`reverse(x)`

Reverse the order of the elements of the list `x`. Atomic types are unaffected by this operation.

```
KeRF> reverse 23 78 94
[94, 78, 23]
KeRF> reverse "backwards"
"sdrawkcab"
KeRF> reverse 5
5
KeRF> reverse {a: 23 56, b:0 1}
a:[23, 56], b:[0, 1]
```

The symbol `/` is equivalent to `reverse` when used as a unary operator:

```
KeRF> /"example text"
"txet elpmaxe"
```

9.135 rsum - Running Sum

`rsum(x)`

Calculate a running sum of the elements of the list `x`, from left to right. nans are ignored.

```
KeRF> rsum 1
1
KeRF> rsum 1 2
[1, 3]
KeRF> rsum 1 2 5
[1, 3, 8]
KeRF> rsum 1 2 5 7 8
[1, 3, 8, 15, 23]
KeRF> rsum 1 2 3 nan 4
[1, 3, 6, 6, 10.0]
KeRF> rsum []
[]
```

9.136 run - Run

`run(filename)`

Load and run Kerf source from a file. Equivalent to `load(filename)`.

9.137 **search** - Search

`search(x, y)`

Look for x in y. If found, return the index. If not found, return NAN:

```
KeRF> search(3, 0 3 17 30)
1
KeRF> search(30, 0 3 17 30)
3
KeRF> search(15, 0 3 17 30)
NAN
KeRF> search("F", "AEIOU")
NAN
```

9.138 **seed_prng** - Set random seed

`seed_prng(x)`

Sets random number generator seed to x.

9.139 **send_async** - Send Asynchronous

`send_async(x, y)`

Given a connection handle x, as obtained with `open_socket`, send a string y to a remote Kerf instance and do not wait for a reply. See **Network I/O**.

9.140 **send_sync** - Send Synchronous

`send_sync(x, y)`

Given a connection handle x, as obtained with `open_socket`, send a string y to a remote Kerf instance, waiting for a reply. See **Network I/O**.

9.141 **setminus** - Set Disjunction

`setminus(x, y)`

Equivalent to `except(x, y)`.

9.142 **shell** - Shell Command

`shell(x)`

Execute a string x containing a shell command as if from `/bin/sh -c x`. See **General I/O**.

9.143 shift - Shift

```
shift(n, x)
shift(n, x, fill)
```

Offset the list `x` by `n` positions, filling shifted-in positions with `fill`.

```
KeRF> shift(4, 1 2 3 4 5 6 7, 999)
[999, 999, 999, 999, 1, 2, 3]
KeRF> shift(-1, 1 2 3 4 5 6 7, 999)
[2, 3, 4, 5, 6, 7, 999]
```

If `fill` is not provided, use a type-appropriate null value as generated by `type_null`.

```
KeRF> shift(3, "ABCDE")
"   AB"
KeRF> shift(-3, "ABCDE")
"DE   "
KeRF> shift(2, 1 2 3 4)
[NAN, NAN, 1, 2]
KeRF> shift(2, 1.0 2.0 3.0 4.0)
[nan, nan, 1, 2.0]
```

9.144 shuffle - Shuffle

```
shuffle(x)
```

Randomly permute the elements of the list `x`. Equivalent to `rand(-len(x), x)`.

```
KeRF> shuffle "APPLE"
"PAEPL"
KeRF> shuffle "APPLE"
"LEAPP"
KeRF> shuffle "APPLE"
"LEPAP"
```

9.145 sin - Sine

```
sin(x)
```

Calculate the sine of `x`, expressed in radians. Atomic. The results of `sin` will always be floating point values.

```
KeRF> sin 3.14159 1 -20
[2.65359e-06, 0.841471, -0.912945]
KeRF> asin sin 3.14159 1 -20
[2.65359e-06, 1, -1.15044]
```


9.146 `sinh` - Hyperbolic Sine

`sinh(x)`

Calculate the hyperbolic sine of `x`, expressed in radians. Atomic. The results of `sinh` will always be floating point values.

```
KerF> sinh 3.14159 1 -20
[11.5487, 1.1752, -2.42583e+08]
```

9.147 `sleep` - Sleep

`sleep(x)`

Delay for at least `x` milliseconds and then return `x`.
Note that time spent sleeping will not be considered during timing.

```
> time ./kerf -x "sleep 3000"
3000

real    0m3.070s
user    0m0.002s
sys     0m0.063s

KerF> timing 1;
KerF> sleep 5000;
0 ms
```

9.148 `sort` - Sort

`sort(x)`

Sort the elements of the list `x` in ascending order. Equivalent to `x[ascend x]`.

```
KerF> sort "ALPHABETICAL"
"AAABCEHILLPT"
KerF> sort 27 18 4 9
[4, 9, 18, 27]
KerF> sort unique "how razorback jumping frogs can level six piqued gymnasts"
" abcdefghijklmnopqrstuvwxyz"
```

9.149 `sort_debug` - Sort Debug

`sort_debug(x)`

Return a map containing debugging information about the list or table `x` which is relevant to the performance and internal datapaths of searching and sorting.

```
KeRF> sort_debug range(4)
{attr_sorted:1, is_array:1, is_enum:0, is_actually_sorted_array:1, is_index:0,
 is_index_working:0, is_table:0, is_table_sorted:0}

KeRF> sort_debug "ACED"
{attr_sorted:0, is_array:1, is_enum:0, is_actually_sorted_array:0, is_index:0,
 is_index_working:0, is_table:0, is_table_sorted:0}

KeRF> sort_debug 27
{attr_sorted:1, is_array:0, is_enum:0, is_actually_sorted_array:0, is_index:0,
 is_index_working:0, is_table:0, is_table_sorted:0}

KeRF> sort_debug {{a: 4 2 1}}
{attr_sorted:0, is_array:0, is_enum:0, is_actually_sorted_array:0, is_index:0,
 is_index_working:0, is_table:1, is_table_sorted:0}
```

9.150 `split` - Split List

`split(x, y)`

Subdivide the list `y` at the index/indices given by `x`. Equivalent to `y[xvals part rsum (xkeys y) in x]`.

```
KeRF> split(4 9, "SomeWordsTogether")
["Some", "Words", "Together"]
KeRF> 3 split "ABCDEF"
["ABC", "DEF"]
KeRF> split(0 1 3 5, 11 22 33 44 55)
[[11], [22, 33], [44, 55]]
```

9.151 `sqrt` - Square Root

`sqrt(x)`

Calculate the square root of `x`. Atomic.

```
KeRF> sqrt 2 25 100
[1.41421, 5, 10.0]
```

9.152 stamp_diff - Timestamp Difference

stamp_diff(x, y)

Calculate the difference between the timestamps x and y in nanoseconds.

```
KeRF> t: now(); sleep(10); stamp_diff(now(), t)
10302000
KeRF> t: now(); sleep(10); stamp_diff(t, now())
-10874000
```

9.153 std - Standard Deviation

std(x)

Calculate the standard deviation of the elements of the list x. Equivalent to `sqrt var x`.

```
KeRF> std 4 7 19 2 0 -2
6.87992
KeRF> std {a: 4 1 0}
{a:1.69967}
```

9.154 string - Cast to String

string(x)

Convert the value x to a string. Atomic. If you wish to recursively convert an entire data structure to a string, use `rep`.

```
KeRF> string 990
"990"
KeRF> string 15 9 10
["15", "9", "10"]
KeRF> string {a:4 5}
{a:["4", "5"]}
```

9.155 subtract - Subtract

subtract(x, y)

Equivalent to `minus(x, y)`.

9.156 sum - Sum

sum(x)

Calculate the sum of the elements of the list x.

```
KeRF> sum 4 3 9
16
KeRF> sum 5
5
KeRF> sum []
0
```

9.157 tables - Tables

tables()

Generate a list of the names of all currently loaded tables.

```
KeRF> tables()
[]
KeRF> t: {{a: 1 2 3, b: 4 5 6}};
KeRF> tables()
["t"]
```

9.158 take - Take

take(x, y)

Create a list containing the first x elements of y, looping y as necessary. If x is negative, take backwards from the last to the first. Equivalent to first(x, y).

```
KeRF> take(3, `"A")
"AAA"
KeRF> take(2, range(5))
[0, 1]
KeRF> take(8, range(5))
[0, 1, 2, 3, 4, 0, 1, 2]
KeRF> take(-3, range(5))
[2, 3, 4]
```

The symbol ^ is equivalent to take when used as a binary operator:

```
KeRF> 3^"ABCDE"
"ABC"
```

9.159 tan - Tangent

tan(x)

Calculate the tangent of x, expressed in radians. Atomic. The results of tan will always be floating point values.

```
KeRF> tan 0.5 -0.2 1 4
[0.546302, -0.20271, 1.55741, 1.15782]
KeRF> atan(tan 0.5 -0.2 1 4)
[0.5, -0.2, 1, 0.858407]
```

9.160 tanh - Hyperbolic Tangent

tanh(x)

Calculate the hyperbolic tangent of x, expressed in radians. Atomic. The results of tanh will always be floating point values.

```
KeRF> tanh 3.14159 1 -20
[0.996272, 0.761594, -1.0]
```

9.161 times - Multiplication

times(x, y)

Calculate the product of x and y. Fully atomic.

```
KeRF> times(3, 5)
15
KeRF> times(1 2 3, 5)
[5, 10, 15]
KeRF> times(3, 5 8 9)
[15, 24, 27]
KeRF> times(10 15 3, 8 2 4)
[80, 30, 12]
```

The symbol * is equivalent to times when used as a binary operator:

```
KeRF> 1 2 3*2
[2, 4, 6]
```

9.162 timing - Timing

timing(x)

If x is truthy, enable timing. Otherwise, disable it. Returns a boolean timing status. When timing is active, all operations will print their approximate runtime in milliseconds after completing.

```
KeRF> timing(1)
1
KeRF> sum range exp(2, 24)
140737479966720

203 ms
KeRF> timing(0)
0
```

9.163 tolower - To Lowercase

tolower(x)

Convert a string x to lowercase. Equivalent to floor(x).

9.164 toupper - To Uppercase

toupper(x)

Convert a string x to uppercase. Equivalent to ceil(x).

9.165 transpose - Transpose

transpose(x)

Take the transpose (flip the x and y axes) of a matrix x. Has no effect on atoms or lists of atoms.

```
KeRF> transpose [1 2 3, 4 5 6, 7 8 9]
[[1, 4, 7],
 [2, 5, 8],
 [3, 6, 9]]
KeRF> transpose 1 2 3
[1, 2, 3]
```

Atoms will “spread” as needed to produce a rectangular matrix if the list contains any sublists:

```
KeRF> transpose [2, 3 4 5, 6]
[[2, 3, 6],
 [2, 4, 6],
 [2, 5, 6]]
```

The symbol + is equivalent to transpose when used as a unary operator:

```
KeRF> +[1 2, 3 4]
[[1, 3],
 [2, 4]]
```

9.166 trim - Trim

trim(x)

Remove leading and trailing whitespace from strings. Atomic.

```
KeRF> trim(" some text ")
"some text"
KeRF> trim [" some text ", " another", "\tab\t"]
["some text", "another", "ab"]
```

9.167 type_null - Type Null

type_null(x)

Generate the equivalent type-specific null value for x.

```
KeRF> type_null("ABC")
~" "
KeRF> type_null(1.0)
nan
KeRF> type_null(1)
NAN
KeRF> type_null(now())
00:00:00.000
```

9.168 `uneval` - Uneval

`uneval(x)`

Equivalent to `json_from_kerf(x)`.

9.169 `union` - Set Union

`union(x, y)`

Construct an unsorted list of the unique elements in either `x` or `y`. Equivalent to `distinct join(x, y)`.

```
KeRF> union(2, 4 5)
[2, 4, 5]
KeRF> union([], 2)
[2]
KeRF> union(2 3 4, 1 3 9)
[2, 3, 4, 1, 9]
```

9.170 `unique` - Unique Elements

`unique(x)`

Equivalent to `distinct(x)`.

9.171 `var` - Variance

`var(x)`

Calculate the variance of the elements of a list `x`. Equivalent to `(sum (x - avg x)**2)/count_nonnull x`.

```
KeRF> var []
nan
KeRF> var 4 3 8 2
5.1875
KeRF> sqrt var 4 3 8 2
2.27761
```

9.172 which - Which

`which(x)`

For each index i of the list `x`, produce `x[i]` copies of `i`:

```
KeRF> which 1 2 1 4
[0, 1, 1, 2, 3, 3, 3, 3]
KeRF> which 1 2 3
[0, 1, 1, 2, 2, 2]
KeRF> which 1 1 1
[0, 1, 2]
```

This operation is most often used to retrieve a list of the indices of nonzero elements of a boolean vector:

```
KeRF> which 0 0 1 0 1 1 0 1
[2, 4, 5, 7]
```

The symbol `?` is equivalent to `which` when used as a unary operator:

```
KeRF> ?0 0 1 0 1 1 0 1
[2, 4, 5, 7]
```

9.173 write_csv_from_table - Write CSV From Table

`write_csv_from_table(filename, table)`

Write table to disk as a Comma-Separated Value file called `filename`.
Equivalent to `write_delimited_file_from_table(",", filename, table)`.

9.174 write_delimited_file_from_table - Write Delimited File From Table

`write_delimited_file_from_table(delimiter, filename, table)`

Write table to disk as `filename` using newlines to separate rows and `delimiter` to separate columns. The file will be written with a header row corresponding to the keys of the columns of `table`. Returns the number of bytes written to the file.

```
KeRF> t: {{a: 1 2 3, b:["one", "two", "three"]}};
KeRF> write_delimited_file_from_table("|", "example.psv", t)
23
KeRF> shell("wc -c example.psv")
["      23 example.psv"]
KeRF> shell("cat example.psv")
["a|b", "1|one", "2|two", "3|three"]
```

9.175 write_stripped_to_path - Write Striped File To Path

`write_stripped_to_path(path, x)`

Write `x` to disk at `path` as a striped file. See **Striped Files**.

9.176 `write_text` - Write Text

```
write_text(filename, x)
```

Write the value `x` to a text file `filename`, creating the file as necessary. See **File I/O**.

9.177 `write_to_path` - Write to Path

```
write_to_path(filename, x)
```

Write the value `x` to a binary file `filename`, creating the file as necessary. See **File I/O**.

9.178 `xbar` - XBar

```
xbar(x, y)
```

Equivalent to `x * floor y/x`.

9.179 `xkeys` - Object Keys

```
xkeys(x)
```

Produce a list of keys for a map, table or list `x`.

```
KeRF> xkeys 33 14 9
[0, 1, 2]
KeRF> xkeys {a: 42, b: 49}
["a", "b"]
KeRF> xkeys {{a: 42, b: 49}}
["a", "b"]
```

9.180 `xvals` - Object Values

```
xvals(x)
```

Produce a list of values for a map or table `x`. If `x` is a list, produce a list of valid indices to `x`.

```
KeRF> xvals 33 14 9
[0, 1, 2]
KeRF> xvals {a: 42, b: 49}
[42, 49]
KeRF> xvals {{a: 42, b: 49}}
[[42], [49]]
```

10 Combinator Reference

10.1 converge - Converge

Given a unary function, apply it to a value repeatedly until it does not change or the next iteration will repeat the initial value. Some functional languages refer to this operation as *fixedpoint*.

```
KeRF> {[x] floor x/2} converge 32
0
KeRF> {[x] mod(x+1, 5)} converge 0
4
KeRF> {[x] display x; mod(x+1, 5)} converge 3
3
4
0
1
2
2
```

If a numeric left argument is provided, instead repeatedly apply the function some number of times:

```
KeRF> 3 {[x] floor x/2} converge 32
4
KeRF> 3 {[x] x*2} converge 32
256
KeRF> 5 {[x] join("A", x)} converge "B"
"AAAAAB"
```

Applied to a binary function, converge is equivalent to fold.

10.2 deconverge - Deconverge

deconverge is similar to converge, except it gathers a list of intermediate results.

```
KeRF> {[x] floor x/2} deconverge 32
[32, 16, 8, 4, 2, 1, 0]
KeRF> {[x] mod(x+1, 5)} deconverge 1
[1, 2, 3, 4, 0]
KeRF> 3 {[x] floor x/2} deconverge 32
[32, 16, 8, 4]
KeRF> 3 {[x] x*2} deconverge 32
[32, 64, 128, 256]
```

Applied to a binary function, deconverge is equivalent to unfold.

10.3 fold - Fold

Given a binary function, apply it to pairs of the elements of a list from left to right, carrying the result forward on each step. Some functional languages refer to this operation as `foldl`.

```
KeRF> add fold 1 2 3 4
10
KeRF> {[a,b] join(enlist a, b)} fold 1 2 3 4
[[[1, 2], 3], 4]
```

Note that the function will not be applied if folded over an empty or 1-length list:

```
KeRF> {[a,b] out "nope"; a+b} fold [5]
5
KeRF> {[a,b] out "nope"; a+b} fold 1 2
nope 3
```

It is also possible to supply an initial value for the fold as a left argument:

```
KeRF> 0 {[a,b] join(enlist a, b)} fold 1 2 3
[[[0, 1], 2], 3]
KeRF> 7 {[a,b] out "yep"; a+b} fold 5
yep 12
```

The symbol `\/` is equivalent to `fold`:

```
KeRF> add \/ 1 2 3
6
KeRF> 7 add \/ 5
12
```

Applied to a unary function, `fold` is equivalent to `converge`.

10.4 mapback - Map Back

Given a binary function, `mapback` pairs up each value of a list with its predecessor and applies the function to these values. The first item of the resulting list will be the first item of the original list:

```
KeRF> join mapback 1 2 3
[1,
 [2, 1],
 [3, 2]]
```

A common application of `mapback` is to calculate deltas between successive elements of a list:

```
KeRF> - mapback 5 3 2 9
[5, -2, -1, 7]
```

If a left argument is provided, it will be used as the previous value of the right argument's first value:

```
KeRF> 1 join mapback 2 3 4
[[2, 1],
 [3, 2],
 [4, 3]]
```

The symbol `\~` is equivalent to `mapback`:

```
KeRF> 0 != \~ 1 1 0 1 0 0 0
      [1, 0, 1, 1, 1, 0, 0]
```

10.5 mapcores - Map to Cores

Similar to `mapdown`, except it distributes work to multiple CPU cores (as available) and performs work in parallel. There is some overhead to distributing work and collecting results, so only use `mapcores` after careful profiling.

```
KeRF> time: {[f] s:now(); v:f(); [v, stamp_diff(now(), s)]}];

KeRF> time {[ ] {[x] sleep 100; x*2} mapdown 1 3 7 9}
      [[2, 6, 14, 18], 404506000]

KeRF> time {[ ] {[x] sleep 100; x*2} mapcores 1 3 7 9}
      [[2, 6, 14, 18], 184118000]
```

Every parallel worker is given its own independent environment tree, and writes to global variables are discarded after results are gathered. Side effects to IO devices are performed in an arbitrary order unless otherwise synchronized.

```
KeRF> a: [0];
KeRF> {[x] a[0]:2*x; a[0]} mapdown 1 3 7 9
      [2, 6, 14, 18]
KeRF> a
      [18]

KeRF> a: [0];
KeRF> {[x] a[0]:2*x; a[0]} mapcores 1 3 7 9
      [2, 6, 14, 18]
KeRF> a
      [0]

KeRF> {[x] display x; 2*x} mapcores 1 3 7 9
3
7
9
1
      [2, 6, 14, 18]
```

Uses of `mapcores` can be nested, generally with diminishing returns. The current implementation arbitrarily caps nesting at 2:

```
KeRF> {[x] {[x] 2*x} mapcores range x} mapcores 3 2 5
      [[0, 2, 4],
       [0, 2],
       [0, 2, 4, 6, 8]]

KeRF> {[x] {[x] {[x] 2*x} mapcores range x} mapcores range x} mapcores 2 3
      ...nge x} mapcores range x} mapcores 2 3
      ^
Parallel execution error
```

10.6 mapdown - Map Down

Apply a unary function to every element of a list, yielding a new list of the same size. Some functional programming languages refer to this as simply map. mapdown can be used to achieve a similar effect to how atomic built-in functions naturally “push down” onto the values of lists.

```
KeRF> negate mapdown 2 -5
[-2, 5]
KeRF> {[n] 3*n} mapdown 2 5 9
[6, 15, 27]
```

Given a binary function and a left argument, mapdown pairs up sequential values from two equal-length lists and applies the function to these pairs. Some functional programming languages refer to this as zip, meshing together a pair of lists like the teeth of a zipper:

```
KeRF> 1 2 3 join mapdown 4 5 6
[[1, 4],
 [2, 5],
 [3, 6]]
```

mapdown also works with maps and tables:

```
KeRF> count mapdown {a: 1 2, b: 3 4 5, c: 6}
[2, 3, 1]
KeRF> reverse mapdown {{a: 1 2, b: 3 4}}
[[2, 1], [4, 3]]
```

The symbol \= is equivalent to mapdown:

```
KeRF> {[n] 3*n} \= 2 5 9
[6, 15, 27]
```

10.7 mapleft - Map Left

Given a binary function, apply it to each of the values of the left argument and the right argument, gathering the results in a list.

```
KeRF> 1 2 3 join mapleft 4
[[1, 4],
 [2, 4],
 [3, 4]]
```

The symbol \> is equivalent to mapleft.

10.8 mapright - Map Right

Given a binary function, apply it to a left argument and each of the values of the right argument, gathering the results in a list. mapright, like mapdown, provides a way of “pushing a function down onto” data or overriding existing atomicity:

```
KeRF> 1 join mapright 2 3 4
[[1, 2],
 [1, 3],
 [1, 4]]
```

mapright and mapleft can be used to take the cartesian product of two lists:

```
KeRF> 0 1 2 add 0 1 2
[0, 2, 4]
KeRF> 0 1 2 add mapright 0 1 2
[[0, 1, 2],
 [1, 2, 3],
 [2, 3, 4]]
```

The symbol \< is equivalent to mapright.

10.9 reconverge - Reconverge

Equivalent to unfold.

10.10 reduce - Reduce

Equivalent to fold.

10.11 refold - Refold

Equivalent to unfold.

10.12 rereduce - Re-Reduce

Equivalent to unfold.

10.13 unfold - Unfold

unfold is similar to fold, except it gathers a list of intermediate results. This can often provide a useful way to debug the behavior of fold.

```
KeRF> add unfold 1 2 3 4
[1, 3, 6, 10]
KeRF> 100 add unfold 1 2 3 4
[101, 103, 106, 110]
KeRF> {[a,b] join(enlist a, b)} unfold 1 2 3 4
[1,
 [1, 2],
 [[1, 2], 3],
 [[[1, 2], 3], 4]]
```

The symbol \\ is equivalent to unfold:

```
KeRF> add \\ 1 2 3
[1, 3, 6]
KeRF> 7 add \\ 5
[12]
```

Applied to a unary function, unfold is equivalent to deconverge.

11 Global Reference

11.1 .Argv - Arguments

A list of the arguments provided to Kerf at the command line.

```
>kerf -q foo
File handle status failed during directory check.: No such file or directory
KerF> .Argv
["kerf", "-q", "foo"]
```

11.2 Math

11.2.1 .Math.BILLION - Billion

Constant representing 10^9 .

11.2.2 .Math.E - E

Constant representing Euler's number. 2.7182818284590452353602.

11.2.3 .Math.TAU - Tau

Constant representing 2π . 6.2831853071795864769252.

11.3 Net

11.3.1 .Net.client - Client

During IPC execution, contains a constant representing the current client's unique handle. See **Network I/O**.

11.3.2 .Net.on_close - On Close

If defined, an IPC server will call this single-argument function with a client handle when that client closes its connection. See **Network I/O**.

11.4 Parse

11.4.1 `.Parse.strptime_format` - Time Stamp Format

Specifies the format used for formatting and parsing dates and time stamps from delimited files when the field specifier is Z. Builds directly on the standard C function `strptime()`, as defined in `time.h`.

By default, `%d-%b-%y %H:%M:%S`.

- `%b` Abbreviated month name
- `%B` Full month name
- `%c` Date and time
- `%C` Century 00-99
- `%d` Day of month as decimal number 01-31
- `%e` Day of the month as decimal number 1-31 with leading space for single digit
- `%F` Equivalent to Y-m-d ISO 8601
- `%h` Equivalent to `%b`
- `%H` Hours as a decimal number 00-23
- `%I` Hours as a decimal number 00-12
- `%j` Day of year as decimal number 001-366
- `%m` Month as decimal number 01-12
- `%M` Minute as decimal number 00-59
- `%p` AM/PM indicator, used with `%I`
- `%S` Second as integer 00-61; aka up to two leap seconds
- `%u` Weekday as a decimal number 1-7, Monday is 1
- `%w` Weekday as a decimal number 0-6, Sunday is 0
- `%W` Week of the year as a 00-53 with Monday as first day of week
- `%y` Year without century
- `%Y` Year with century
- `%z` Signed offset in hours and minutes from UTC
- `%Z` Time zone abbreviation as a character string

11.4.2 `.Parse.strptime_format2` - Time Format

Specifies the format used for formatting and parsing timestamps from delimited files when the field specifier is Y. Builds directly on the standard C function `strptime()`, as defined in `time.h`. Used for parsing time-only columns in data files where dates and times are in separate columns.

12 Programming Techniques

This section contains case studies illustrating how Kerf can be used to solve problems, ranging from simple to complex. We will build up solutions step by step and consider tradeoffs in performance and style.

12.1 Reversing a Map

A map links a set of keys with a set of values. It is easy and efficient to use a key to look up the associated value. Sometimes it is desirable to go the other way- using a value to look up the associated key. One way to approach this problem might be to iterate through each key in the map until we found one which is associated with our target value:

```
def find_key(m, val) {
  keys: xkeys m
  for(i: 0; i < len(keys); i: i+1) {
    if (m[keys[i]] == val) { return keys[i] }
  }
  return null
}
```

As the number of keys in the map scales up, the amount of time this function takes to run increases proportionally. If we need to perform many repeated reverse lookups, this approach will be prohibitively slow.

As a general rule of thumb when programming, if something is too expensive to calculate, cache it. By writing a function which analyzes a map and constructs a new map which “reverses” the keys and values of the original, we can pay this construction cost once and then achieve efficient reverse-lookups.

There are several ways to approach this problem. A programmer used to imperative programming languages might come up with something like this:

```
def map_reverse_1(m) {
  r: {}
  keys: xkeys m
  vals: xvals m
  for(i: 0; i < len(keys); i: i+1) {
    r[vals[i]]: keys[i]
  }
  return r
}
```

Perhaps you want to get fancy and try to represent the same algorithm using the combinator fold?

```
{ } { [m, e] m[e[0]]: e[1] } fold transpose([xvals m, xkeys m])
```

The simplest approach is to take advantage of the map built-in function:

```
def map_reverse_2(m) {
  return map(xvals m, xkeys m)
}
```

All of these solutions are making an unsafe assumption. The keys of a map are always unique by definition, but what happens if the values are not unique?

```
KeRF> map_reverse_2({a: "A", b:"B", c:"C"})
{A:"a", B:"b", C:"c"}
KeRF> map_reverse_2({a: "A", b:"B", c:"A"})
{A:"a", B:"b"}
```

We've lost information! A more robust map reversal routine should gather the keys of repeated values as a list, producing a result more like:

```
KeRF> map_reverse_3({a: "A", b:"B", c:"C"})
{A:["a","c"], B:["b"]}
```

First, let's look at an imperative approach to solving the problem. We can iterate through the keys of the original map, as before. Instead of storing the key directly in the result map, we will append it to a list stored in the map. This requires us to first ensure that any slot in the result map is initialized with an empty list:

```
def map_reverse_3(m) {
  r:
  keys: xkeys m
  vals: xvals m
  for(i: 0; i < len(keys); i: i+1) {
    if (not r has_key vals[i]) { r[vals[i]]: [] }
    r[vals[i]]: join(r[vals[i]], enlist keys[i])
  }
  return r
}
```

The enlist is important here- without it, string keys would be mashed together:

```
KeRF> join([], "foo")
"foo"
KeRF> join("foo", "bar")
"foobar"
KeRF> join([], enlist "foo")
["foo"]
KeRF> join(["foo"], enlist "bar")
["foo", "bar"]
```

We can simplify this slightly by using “spread assignment” to initialize our result map with empty lists in every entry:

```
...
r[vals]: [[]]
for(i: 0; i < len(keys); i: i+1) {
  r[vals[i]]: join(r[vals[i]], enlist keys[i])
}
...
```

What happens if we leave out this initialization entirely? Try it!

Now let's look at a functional approach to solving this problem. If we partition the value set of the original map, we obtain lists of the indices of the values which are identical:

```
KerF> m: {a: "A", b:"B", c:"A"}
      {a:"A", b:"B", c:"A"}
KerF> xvals m
["A", "B", "A"]
KerF> part xvals m
{A:[0, 2], B:[1]}
```

Since the key and value vectors produced by `xkeys` and `xvals` line up, we could also use those indices to obtain the corresponding keys for each group of values. Kerf's powerful indexing facilities make this simple:

```
KerF> xkeys m
["a", "b", "c"]
KerF> xvals part xvals m
[[0, 2], [1]]
KerF> (xkeys m)[xvals part xvals m]
[["a", "c"], ["b"]]
```

Now we have what we need to assemble the result, so we use `map` to join the unique elements of the original value set with the lists of corresponding keys we computed previously. The `unique` function collects results in the order they appear, just like `part`, so everything will line up properly:

```
def map_reverse_4(m) {
    nk: unique xvals m
    nv: (xkeys m)[xvals part xvals m]
    return map(nk, nv)
}
```

Notice how we were able to use the Kerf REPL and a single working example to experiment interactively and then distill the results into a general definition. It is very natural to develop functional solutions to problems in this manner. Manipulating an entire data structure at once often results in a simpler solution with fewer conditionals and loops than trying to solve a problem “one step at a time”.

It is also possible to write this as a more compact one-liner by avoiding intermediate variables and using some of the symbolic shorthands for `map` (!), `unique` (%), `enumerate` (^) and `part` (&), but the result is much harder to read and understand at a glance:

```
{[m] (%xvals m)!(^m)[xvals(&xvals m)]}
```

Shorthand is a nice way to save typing at the REPL, but favor a more relaxed, verbose style when writing larger programs- future maintainers (yourself included) will be thankful! Use Kerf's syntactic alternatives- infix function calls, optional parentheses, function aliases, etc.- to write your programs in the clearest, most readable way possible.

12.2 Run-Length Encoding

Run-Length Encoding (RLE) is a simple means of compressing data. Wherever there is a repeated run of identical elements, instead of storing each element store a single copy of the element and the length of the run.

For example, if our input list looked like:

```
"AAABBAAAAAAAA"
```

An RLE-compressed representation might look like:

```
[[3, `A"], [2, `B"], [7, `A"]]
```

Let's start by writing a *decompressor*- given an RLE-compressed list, reconstruct the original. For a given element of the list, we can use `take` to create a run:

```
KeRF> p: [3, `A"]
      [3, `A"]
KeRF> take(p[0], p[1])
      "AAA"
```

One way to apply a binary function to a pair of arguments is to use the combinator `fold`. A fold is like placing a function between the elements of a list. Thus, these two expressions are equivalent:

```
KeRF> take fold p
      "AAA"
KeRF> p[0] take p[1]
      "AAA"
```

We really want to apply `take` to *each* pair in the original list. The combinator `mapdown` applies a unary function to each element of a list, returning a list of results:

```
KeRF> (take fold) mapdown [[3, `A],[2, `B],[7, `A]]
      ["AAA", "BB", "AAAAAAAA"]
```

The parentheses are not required- in this case they are simply making it clearer that the function `mapdown` is applying to each element of the list to the right is “`take fold`”. When combinators are attached to functions, they form new compound functions which can be passed around or stored in variables just like any other function:

```
KeRF> take fold mapdown [[3, `A],[2, `B],[7, `A]]
      ["AAA", "BB", "AAAAAAAA"]
KeRF> take fold
      take fold
```

To reproduce the original sequence from this list of runs, we need to join all the runs together. The built-in function `flatten` does exactly what we need. Alternatively, we could use the equivalent `join fold`:

```
KeRF> flatten ["AAA", "BB", "AAAAAAAA"]
      "AAABBAAAAAAAA"
KeRF> join fold ["AAA", "BB", "AAAAAAAA"]
      "AAABBAAAAAAAA"
```

Putting everything together, we could write a definition for `rle_decode` in several ways:

```
def rle_decode(x) {return flatten take fold mapdown x}

rle_decode: {[x] flatten take fold mapdown x}

rle_decode: flatten take fold mapdown
```

This last approach is called a *tacit definition*- it doesn't require naming any variables or arguments. The combinators glue together functions to produce a function that is simply waiting for a right argument. Composing together functions won't always work out this nicely, but when a tacit definition is possible the results can be very aesthetically pleasing, as if there were hardly any syntax to the language at all:

```
KeRF> rle_decode: flatten take fold mapdown
      flatten take fold mapdown

KeRF> rle_decode [[3, `A],[2, `B],[7, `A]]
      "AAABBAAAAAAAA"

KeRF> flatten take fold mapdown [[3, `A],[2, `B],[7, `A]]
      "AAABBAAAAAAAA"
```

Now that we're a bit more comfortable with using combinators, let's tackle the *compressor*- given a list, we want to break it into runs. For each run, we need to determine the first element (or any element, really) and the length of the run.

To find the start of each run, we can compare each element of the list with the item which came before it. If they're different, we are beginning a new run. Otherwise, we're continuing an existing run. The combinator `mapback` applies a function to each element of a list and its predecessor, which is just what we're looking for:

```
KeRF> != mapback "AAABBAAAAAAAA"
      [ `A", 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0]
```

Well, *nearly* what we're looking for. What's that character doing at the beginning of our result? `mapback` doesn't have anything to pair the first element of a list up with, so by default it leaves that item alone. If we supply a value on the left, `mapback` will use that to pair with the first element of the list on the right. Supplying a null will ensure that the first list element is considered "not equal to" this default value:

```
KeRF> `A != mapback "AAABBAAAAAAAA"
      [0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0]

KeRF> `B != mapback "AAABBAAAAAAAA"
      [1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0]

KeRF> `B != mapback "BAABBAAAAAAAA"
      [0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0]

KeRF> null != mapback "AAABBAAAAAAAA"
      [1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0]
```

The function which, given a boolean list, produces a list of the indices of the 1s. If we index into the original list, we can retrieve a list of the items which begin each run, in order:

```
KeRF> s: "AAABBAAAAAAAA"
      "AAABBAAAAAAAA"
KeRF> null != mapback s
      [1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0]
KeRF> which null != mapback s
      [0, 3, 5]
KeRF> s[which null != mapback s]
      "ABA"
```

Now we just need to find the lengths of each run. Let's consider that vector we already created which identifies run heads. If we take a running sum (rsum) of that list, we can uniquely label all the members of each run:

```
KeRF> h: null != mapback s
      [1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0]
KeRF> rsum h
      [1, 1, 1, 2, 2, 3, 3, 3, 3, 3, 3, 3]
```

Partitioning these labels will conveniently group them together, and we can count each of the resulting groups:

```
KeRF> part rsum h
      1:[0, 1, 2], 2:[3, 4], 3:[5, 6, 7, 8, 9, 10, 11]
KeRF> count mapdown part rsum h
      [3, 2, 7]
```

With the list of counts and the list of values, we can join each to produce the tuples we want. By giving the combinator mapdown a left argument and a binary function, we can neatly “zip” together our lists:

```
KeRF> 1 2 3 join mapdown "ABC"
      [[1, `A"],
       [2, `B"],
       [3, `C"]]
```

Bringing everything together, we can arrive at a definition like this:

```
def rle_encode(s) {
  heads:  null != mapback s
  vals:   s[which heads]
  counts: count mapdown part rsum heads
  return  counts join mapdown vals
}
```

For the sake of comparison, here's an imperative solution:

```
def rle_encode_2(s) {
  r: []
  c: 1
  for(i: 0; i < len(s); i: i+1) {
    if (s[i] != s[i+1]) {
      r: join(r, enlist [c, s[i]])
      c: 1
    } else {
      c: c+1
    }
  }
  return r
}
```

Which of these is easier to understand and reason about? The imperative program is familiar, but involves a number of variables which continuously change throughout execution. Did we make any off-by-one errors? The functional solution replaces loops and conditionals with built-in functions and combinators, and never changes the contents of a variable once it is assigned.

“Readability” is inherently subjective, so let's compare performance:

```
KeRF> data: 10000?"ABCD";
KeRF> timing 1;
KeRF> rle_encode data ;
      21 ms
KeRF> rle_encode_2 data ;
      781 ms
```

Kerf executes the functional solution substantially faster. Combinators and built-in functions which operate on entire lists at once are highly optimized and have very little interpreter overhead, so an individual operation can approach native execution speeds. Explicit loops and conditionals force the interpreter to do most of the heavy lifting, which is less efficient.

12.3 HTTP Fetching

KeRF is designed for processing and analyzing data. Many websites exist today which expose access to interesting datasets via an HTTP (HyperText Transfer Protocol) interface. Let's look at HTTP and the performance implications of different styles of data import.

We will use <https://www.quandl.com> for our examples. Quandl provides a wide variety of financial and economic datasets, many of which are available free of charge. Creating an account with Quandl and using credentials with API requests increases the number of queries you are allowed to perform daily, but the system is usable without any form of registration.

In an HTTP transaction, a client makes a *request* for a resource, and the server provides a *response*. Requests have an associated *verb*, which indicates what the client wishes the server to do. The most basic type of request verb is GET- a request for the retrieval of a named resource, indicated by a URL. The response will contain a *response code* indicating success or a class of failure which occurred, and sometimes a *payload*. HTTP is stateless: each request-response pair is an independent transaction, and requests must include all context necessary for processing.

The simplest way to perform an HTTP request is via the Unix curl utility. By using the shell function, we can invoke curl from Kerf. In this example, using the .json suffix on the URL instructs Quandl to produce a response in the JSON format, which we can easily parse. shell breaks input on newlines, so it is necessary to join these back together via implode. The -s option prevents curl from printing information about its download progress, which is not relevant for us.

```
KeRF> url: "https://www.quandl.com/api/v3/datasets/WIKI/AAPL/data.json?rows=10";  
KeRF> d: kerf_from_json implode(`\n`, shell "curl -s "#url#" -X GET");
```

It may also be a good idea to use the -o flag option with curl, which makes it write the response to a temporary file for later retrieval:

```
>curl -s -o tempfile.dat https://url.com/path -X GET
```

The response contains a great deal of potentially useful metadata. We can obtain column names:

```
KeRF> d["dataset_data"]["column_names"]  
["Date", "Open", "High", "Low", "Close", "Volume", "Ex-Dividend", "Split Ratio", "Adj. Open",  
"Adj. High", "Adj. Low", "Adj. Close", "Adj. Volume"]
```

The data itself is represented in a row-oriented fashion:

```
KeRF> 2 first d["dataset_data"]["data"]  
[["2015-12-08", 117.52, 118.6, 116.86, 118.23, 34086875.0, 0.0, 1.0, 117.52, 118.6, 116.86,  
118.23, 34086875.0],  
["2015-12-07", 118.98, 119.86, 117.81, 118.28, 31801965.0, 0.0, 1.0, 118.98, 119.86, 117.81,  
118.28, 31801965.0]]
```


To build a table, we can combine the column headings with the transpose of the row data:

```
KeRF> d_data: d["dataset_data"]["data"];
KeRF> d_cols: d["dataset_data"]["column_names"];
KeRF> t: INSERT INTO {{{}} VALUES map(d_cols, transpose d_data);
KeRF> SELECT Date, Open, High, Low FROM t
```

Date	Open	High	Low
2015-12-08	117.52	118.6	116.86
2015-12-07	118.98	119.86	117.81
2015-12-04	115.29	119.25	115.11
2015-12-03	116.55	116.79	114.22
2015-12-02	117.05	118.11	116.08
2015-12-01	118.75	118.81	116.86
2015-11-30	117.99	119.41	117.75
2015-11-27	118.29	118.41	117.6
...

Bringing everything together, we can write a simple helper function for fetching Quandl datasets and loading them into a table, permitting any kind of further querying we like:

```
def quandl_table(query) {
  url: "https://www.quandl.com/api/v3/#query"
  resp: kerf_from_json implode(`"\n", shell "curl -s "#url#" -X GET")
  cols: resp["dataset_data"]["column_names"]
  data: resp["dataset_data"]["data"]
  return INSERT INTO {{{}} VALUES map(cols, transpose data)
}
```

```
KeRF> SELECT Date, Volume FROM quandl_table "datasets/WIKI/AAPL/data.json?rows=10"
```

Date	Volume
2015-12-08	34086875.0
2015-12-07	31801965.0
2015-12-04	56351301.0
2015-12-03	40935107.0
2015-12-02	32793916.0
2015-12-01	34501246.0
2015-11-30	37074611.0
2015-11-27	13038955.0
...	...

If we load the entire dataset instead of limiting the result to the first 10 rows, we get 8824 rows and the payload is roughly 1.2mb. Our `quandl_table` routine takes somewhere around 3 seconds to run on an average laptop. Let's dig in and try to identify any performance bottlenecks by using Kerf's nanosecond-accuracy timer:

```
def itemize(f) {
  start: now()
  ticks: f({})
  timespans: map(xkeys ticks, start stamp_diff mapback ticks)
  display timespans
  display sum xvals timespans
  display timespans / (sum xvals timespans)
}

def benchmark_json(ticks) {
  d: shell "curl -s https://www.quandl.com/api/v3/datasets/WIKI/AAPL/data.json -X GET"
  ticks["curl"]: now()

  j: kerf_from_json implode(`"\n", d)
  ticks["parse"]: now()

  cols: j["dataset_data"]["column_names"]
  data: j["dataset_data"]["data"]
  t: INSERT INTO {{{}} VALUES map(cols, transpose data);
  ticks["build"]: now()

  return ticks
}
```

```
KerF> itemize(benchmark_json)
curl:5463299000, parse:3175040000, build:6035000
8644374000
curl:0.632006, parse:0.367296, build:0.000698142
```

As you can see, the time spent in `shell` and `curl` is dominant, and represents room for improvement. Parsing the JSON also takes a significant amount of time, and `INSERT INTO` is nearly instantaneous. Quandl can also emit data in a CSV format. How does this compare?

```
def benchmark_csv(ticks) {
  shell "curl -s -o t.csv https://www.quandl.com/api/v3/datasets/WIKI/AAPL/data.csv -X GET"
  ticks["curl"]: now()

  t: read_table_from_csv("t.csv", "SFFFFFFFFFFFF", 1)
  ticks["build"]: now()

  return ticks
}
```

```
KerF> itemize(benchmark_csv)
curl:7386948000, build:73101000
7460049000
curl:0.990201, build:0.009799
```

In this run, we're about 15 percent faster overall, with nearly all our time spent in `shell` and `curl`. This isn't surprising, as CSV is a much simpler file format with less structure than JSON. Clearly, if you're working with large amounts of tabular data, CSV files are more efficient.

Is there another approach? Let's try writing a dynamic library in C which extends Kerf with the IO capabilities we want. By leveraging `libcurl`, we should be able to write a simple routine which performs an HTTP GET for a specified resource and parses the result in JSON. By removing the overhead of interacting with `stdio` or the filesystem it may be possible to improve performance over using the `curl` command-line utility. For general information about dynamic libraries in Kerf, see **FFI**.

For this project we will need to work with Kerf strings. A KERF structure represents strings with a type of `KERF_CHARVEC`, and stores an *unterminated* string in the "g" field. Here's a C function which takes a Kerf string and prints out some information about it:

```
KERF string_info(KERF str) {
    printf("type was %d\n",      str->t);
    printf("length was %d\n",    (int)str->n);
    printf("first char was %c\n", str->g[0]);
    printf("string was '%s'\n",  (int)query->n, query->g);
    return 0;
}
```

For the sake of brevity, examples here will contain a minimum of error checking. They are intended to convey concepts rather than to be ideal production code. In practice when you write C functions to use with Kerf you will want to add code to verify the types of arguments you accept and print meaningful error messages if the function is invoked incorrectly.

We can produce our own strings by using the api function `kerf_api_new_charvec()`, but this copies a complete C string into a KERF structure. For more control, you can instead use `kerf_api_new_kerf()` to allocate an appropriately typed and sized buffer and write into it at your convenience.

On the following page is a complete C program which can be called from Kerf to perform an HTTP GET of an arbitrary URL. For the sake of simplicity, it assumes that responses will be less than 2mb and preallocates space in a KERF structure for this response text.

Ensure that the `kerf_api.h` header file, which should be included along with the Kerf binary, is accessible, and install `libcurl` if necessary. Compile the example as follows:

```
>cc -m64 -flat_namespace -undefined suppress -dynamiclib curly.c -lcurl -o curly.dylib
```

The `-m64` option requests code be generated for a 64-bit architecture. The `-flat_namespace` and `-undefined suppress` flags are present for accessing Kerf's dylib API, as the Kerf binary itself supplies those symbols. The `-lcurl` flag instructs the linker to reference `libcurl`.

```

#include <stdlib.h>
#include <string.h>
#include <curl/curl.h>
#include "kerf_api.h"

#define MAX_PAYLOAD 2097152 // 2mb

static size_t write_data(void *data, size_t size, size_t nmemb, void *destination) {
    KERF payload = (KERF)destination;
    size_t towrite = size*nmemb;
    // abort if the payload is larger than our preallocated buffer
    if ((payload->n) + towrite >= MAX_PAYLOAD) { return 0; }
    memcpy((payload->g) + (payload->n), data, towrite);
    payload->n += towrite;
    return towrite;
}

KERF http_get(KERF url) {
    // copy provided url into a null-terminated c string
    char* url_terminated = malloc((url->n)+1);
    strncpy(url_terminated, url->g, url->n);
    url_terminated[url->n] = '\0';

    // preallocate a fixed size buffer for the result
    KERF payload = kerf_api_new_kerf(KERF_CHARVEC, MAX_PAYLOAD);
    payload->n = 0;

    // instruct libcurl to fetch the URL and block for it
    CURL *curl_handle;
    curl_global_init(CURL_GLOBAL_ALL);
    curl_handle = curl_easy_init();
    curl_easy_setopt(curl_handle, CURLOPT_URL, url_terminated);
    curl_easy_setopt(curl_handle, CURLOPT_NOPROGRESS, 1L);
    curl_easy_setopt(curl_handle, CURLOPT_WRITEFUNCTION, write_data);
    curl_easy_setopt(curl_handle, CURLOPT_WRITEDATA, payload);
    curl_easy_perform(curl_handle);
    curl_easy_cleanup(curl_handle);

    return payload;
}

```

After all that work, let's see how our results compare with the previous approaches:

```
http_get: dlopen("curly.dylib", "http_get", 1)

def benchmark_raw(ticks) {
  d: http_get("https://www.quandl.com/api/v3/datasets/WIKI/AAPL/data.json")
  ticks["curl"]: now()

  j: kerf_from_json d
  ticks["parse"]: now()

  cols: j["dataset_data"]["column_names"]
  data: j["dataset_data"]["data"]
  t: INSERT INTO VALUES map(cols, transpose data);
  ticks["build"]: now()

  return ticks
}
```

```
KeRF> itemize(benchmark_raw)
curl:6428259000, parse:3004744000, build:5596000
9438599000
curl:0.681061, parse:0.318346, build:0.000592885
```

It's actually *slower*! Fetching data over a network will vary wildly in performance, and there simply isn't much we can control here. There's a valuable lesson here, though- by profiling before trying alternatives, we maintain a clear view of what our "optimizations" are getting us. In this case, it seems that dropping into C doesn't pay off, but now we better understand how to do it when necessary. The only way to know for sure whether this approach is appropriate for your application is to try it and take measurements.